



Guide to OpenIG

Version 2.1.0

Paul Bryan
Mark Craig
Jamie Nelson

ForgeRock AS
33 New Montgomery St.,
Suite 1500
San Francisco, CA 94105, USA
+1 415-523-0772
www.forgerock.com

Copyright © 2011-2014 ForgeRock AS



This work is licensed under the [Creative Commons Attribution-NonCommercial-NoDerivs 3.0 Unported License](http://creativecommons.org/licenses/by-nc-nd/3.0/).

To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc-nd/3.0/> or send a letter to Creative Commons, 444 Castro Street, Suite 900, Mountain View, California, 94041, USA.

ForgeRock™ is the trademark of ForgeRock Inc. or its subsidiaries in the U.S. and in other countries. Trademarks are the property of their respective owners.

UNLESS OTHERWISE MUTUALLY AGREED BY THE PARTIES IN WRITING, LICENSOR OFFERS THE WORK AS-IS AND MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND CONCERNING THE WORK, EXPRESS, IMPLIED, STATUTORY OR OTHERWISE, INCLUDING, WITHOUT LIMITATION, WARRANTIES OF TITLE, MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, NONINFRINGEMENT, OR THE ABSENCE OF LATENT OR OTHER DEFECTS, ACCURACY, OR THE PRESENCE OF ABSENCE OF ERRORS, WHETHER OR NOT DISCOVERABLE. SOME JURISDICTIONS DO NOT ALLOW THE EXCLUSION OF IMPLIED WARRANTIES, SO SUCH EXCLUSION MAY NOT APPLY TO YOU.

EXCEPT TO THE EXTENT REQUIRED BY APPLICABLE LAW, IN NO EVENT WILL LICENSOR BE LIABLE TO YOU ON ANY LEGAL THEORY FOR ANY SPECIAL, INCIDENTAL, CONSEQUENTIAL, PUNITIVE OR EXEMPLARY DAMAGES ARISING OUT OF THIS LICENSE OR THE USE OF THE WORK, EVEN IF LICENSOR HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

DejaVu Fonts

Bitstream Vera Fonts Copyright

Copyright (c) 2003 by Bitstream, Inc. All Rights Reserved. Bitstream Vera is a trademark of Bitstream, Inc.

Permission is hereby granted, free of charge, to any person obtaining a copy of the fonts accompanying this license ("Fonts") and associated documentation files (the "Font Software"), to reproduce and distribute the Font Software, including without limitation the rights to use, copy, merge, publish, distribute, and/or sell copies of the Font Software, and to permit persons to whom the Font Software is furnished to do so, subject to the following conditions:

The above copyright and trademark notices and this permission notice shall be included in all copies of one or more of the Font Software typefaces.

The Font Software may be modified, altered, or added to, and in particular the designs of glyphs or characters in the Fonts may be modified and additional glyphs or characters may be added to the Fonts, only if the fonts are renamed to names not containing either the words "Bitstream" or the word "Vera".

This License becomes null and void to the extent applicable to Fonts or Font Software that has been modified and is distributed under the "Bitstream Vera" names.

The Font Software may be sold as part of a larger software package but no copy of one or more of the Font Software typefaces may be sold by itself.

THE FONT SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT OF COPYRIGHT, PATENT, TRADEMARK, OR OTHER RIGHT. IN NO EVENT SHALL BITSTREAM OR THE GNOME FOUNDATION BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, INCLUDING ANY GENERAL, SPECIAL, INDIRECT, INCIDENTAL, OR CONSEQUENTIAL DAMAGES, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF THE USE OR INABILITY TO USE THE FONT SOFTWARE OR FROM OTHER DEALINGS IN THE FONT SOFTWARE.

Except as contained in this notice, the names of Gnome, the Gnome Foundation, and Bitstream Inc., shall not be used in advertising or otherwise to promote the sale, use or other dealings in this Font Software without prior written authorization from the Gnome Foundation or Bitstream Inc., respectively. For further information, contact: fonts@gnome.org.

Arev Fonts Copyright

Copyright (c) 2006 by Tavmjong Bah. All Rights Reserved.

Permission is hereby granted, free of charge, to any person obtaining a copy of the fonts accompanying this license ("Fonts") and associated documentation files (the "Font Software"), to reproduce and distribute the modifications to the Bitstream Vera Font Software, including without limitation the rights to use, copy, merge, publish, distribute, and/or sell copies of the Font Software, and to permit persons to whom the Font Software is furnished to do so, subject to the following conditions:

The above copyright and trademark notices and this permission notice shall be included in all copies of one or more of the Font Software typefaces.

The Font Software may be modified, altered, or added to, and in particular the designs of glyphs or characters in the Fonts may be modified and additional glyphs or characters may be added to the Fonts, only if the fonts are renamed to names not containing either the words "Tavmjong Bah" or the word "Arev".

This License becomes null and void to the extent applicable to Fonts or Font Software that has been modified and is distributed under the "Tavmjong Bah Arev" names.

The Font Software may be sold as part of a larger software package but no copy of one or more of the Font Software typefaces may be sold by itself.

THE FONT SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT OF COPYRIGHT, PATENT, TRADEMARK, OR OTHER RIGHT. IN NO EVENT SHALL TAVMJONG BAH BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, INCLUDING ANY GENERAL, SPECIAL, INDIRECT, INCIDENTAL, OR CONSEQUENTIAL DAMAGES, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF THE USE OR INABILITY TO USE THE FONT SOFTWARE OR FROM OTHER DEALINGS IN THE FONT SOFTWARE.

Except as contained in this notice, the name of Tavmjong Bah shall not be used in advertising or otherwise to promote the sale, use or other dealings in this Font Software without prior written authorization from Tavmjong Bah. For further information, contact: tavmjong@free.fr.

Table of Contents

Preface	v
1. Who Should Use this Guide	v
2. Formatting Conventions	v
3. Accessing Documentation Online	vi
4. Joining the ForgeRock Community	vii
1. Solutions Overview	1
1.1. Extend SSO To Any Application	1
1.2. Federate Enabling Applications	2
1.3. Implement Standards Based Policy Enforcement	2
2. How OpenIG Works	5
2.1. How It Works In A Nutshell	5
2.2. How It Works In Detail	5
3. Getting Started	9
3.1. Before You Begin	9
3.2. Installation Steps	9
3.3. Configuration Steps	11
3.4. Try It Out	11
4. Installing OpenIG	13
5. Configuring Deployment Containers	15
5.1. Tomcat	15
5.2. Jetty	17
5.3. JBoss	18
6. Detailed Use Cases	21
6.1. Portal Application Login	21
6.2. OpenAM Integration	22
6.3. Federation Gateway SP Initiated SAML2 SSO	23
6.4. Federation Gateway IDP Initiated SAML2 SSO	25
6.5. Standards Based SAML2 Agent	25
7. Tutorial With WordPress	27
7.1. Detailed Flow	27
7.2. Setup Summary	28
7.3. Installing OpenIG	29
7.4. Configuring the Browser Host	29
7.5. Configuring OpenIG & Trying It Out	29
7.6. Login with Hard Coded Credentials	30
7.7. Login with Credentials From a Flat File	31
7.8. Login with Credentials From MySQL	32
8. Tutorial On OpenAM Password Capture & Replay	33
8.1. Detailed Flow	33
8.2. Setup Summary	35
8.3. Setup Details	35
9. Using the Federation Gateway	39
9.1. Installation Overview	39

9.2. Configuration File Overview	40
9.3. Configuring the Federation Service	41
9.4. Config.json Sample	41
9.5. Example Settings	43
9.6. Identity Provider Metadata	44
10. Tutorial For the Federation Gateway	45
10.1. Before You Start	45
10.2. Configuring OpenAM	45
10.3. Configuring OpenIG For Federation	46
10.4. Trying It Out	46
11. Configuration Templates	47
11.1. Proxy & Capture	47
11.2. Simple Login Form	49
11.3. Login Form With Cookie From Login Page	50
11.4. Login Form With Extract Filter & Cookie Filter	52
11.5. Login Which Requires a Hidden Value From the Login Page	54
11.6. HTTP & HTTPS Application	56
11.7. Multiple Applications	57
11.8. OpenAM Integration With Headers	59
11.9. Microsoft Online Outlook Web Access	60
12. Customizing OpenIG	63
12.1. Key Extension Points	63
12.2. Implementing a Filter	63
12.3. Implementing a Handler	64
12.4. Heap Object Configuration	64
12.5. Including Servlets and Servlet Filters	65
12.6. Sample Filter	65
13. Troubleshooting	67
13.1. Object not found in heap	67
13.2. Could not find local configuration file	67
13.3. Unexpected character (x) at position 1103	68
13.4. The values in the flat file are incorrect	68
13.5. Problem accessing URL	68
13.6. StaticResponseHandler results in a blank page	69
13.7. OpenIG is not logging users in	69
Index	71

Preface

This guide shows you how to install and configure OpenIG, a high-performance reverse proxy server with specialized session management and credential replay functionality.

1 Who Should Use this Guide

This guide is written for access management designers and administrators who build, deploy, and maintain identity gateways for their organizations. This guide covers the tasks you might perform once or repeat throughout the life cycle of an OpenIG release used in your organization.

You do not need to be an OpenIG wizard to learn something from this guide, though a background in access management and maintaining web application software can help. You do need some background in managing services on your operating systems and in your application servers. You can nevertheless get started with this guide, and then learn more as you go along.

2 Formatting Conventions

Most examples in the documentation are created on GNU/Linux or Mac OS X. Where it is helpful to make a distinction between operating environments, examples for UNIX, GNU/Linux, Mac OS X, and so forth are labeled (UNIX). Mac OS X specific examples can be labeled (Mac OS X). Examples for Microsoft Windows can be labeled (Windows). To avoid repetition, however, file system directory names are often given only in UNIX format as in `/path/to/server`, even if the text applies to `C:\path\to\server` as well.

Absolute path names usually begin with the placeholder `/path/to/`. This path might translate to `/opt/`, `C:\Program Files\`, or somewhere else on your system.

Command line, terminal sessions are formatted as follows.

```
$ echo $JAVA_HOME
/path/to/jdk
```

Command output is sometimes formatted for narrower, more readable output even though formatting parameters are not shown in the command. In the following example, the query string parameter `_prettyPrint=true` is omitted.

```
$ curl https://bjensen:hifalutin@opendj.example.com:8443/users/newuser
{
  "_rev" : "000000005b337348",
  "schemas" : [ "urn:scim:schemas:core:1.0" ],
  "contactInformation" : {
    "telephoneNumber" : "+1 408 555 1212",
    "emailAddress" : "newuser@example.com"
  },
  "_id" : "newuser",
  "name" : {
    "familyName" : "New",
    "givenName" : "User"
  },
  "userName" : "newuser@example.com",
  "displayName" : "New User",
  "meta" : {
    "created" : "2014-06-03T09:58:27Z"
  },
  "manager" : [ {
    "_id" : "kvaughan",
    "displayName" : "Kirsten Vaughan"
  } ]
}
```

Program listings are formatted as follows.

```
class Test {
    public static void main(String [] args) {
        System.out.println("This is a program listing.");
    }
}
```

3 Accessing Documentation Online

ForgeRock core documentation, such as what you are now reading, aims to be technically accurate and complete with respect to the software documented.

Core documentation therefore follows a three-phase review process designed to eliminate errors.

- Product managers and software architects review project documentation design with respect to the users' software lifecycle needs.
- Subject matter experts review proposed documentation changes for technical accuracy and completeness with respect to the corresponding software.
- Quality experts validate implemented documentation changes for technical validity with respect to the software, technical completeness with respect to the scope of the document, and usability for the expected audience.

The review process helps to ensure that documentation published for a ForgeRock release is technically accurate and complete.

Fully reviewed, published core documentation is available at <http://docs.forgerock.org/>. Use this documentation when working with a ForgeRock Enterprise release.

In-progress documentation can be found at each project site under the [Developer Community](#) projects page. Use this documentation when trying a nightly build.

The ForgeRock [Community Wikis](#) provide additional, user-created information. We encourage you to [join the community](#), so that you can update the Wikis, too.

4 Joining the ForgeRock Community

After you [sign up](#) to join the ForgeRock community, you can edit the [Community Wikis](#), and also log bugs and feature requests in the [issue tracker](#).

If you have a question regarding a project but cannot find an answer in the project documentation or Wiki, browse to the [Developer Community](#) page for the project, where you can find details on joining the project mailing lists, and find links to mailing list archives. You can also suggest updates to documentation through the [ForgeRock docs mailing list](#).

The Community Wikis describe how to check out and build source code. Should you want to contribute a patch, test, or feature, or want to author part of the core documentation, first have a look on the ForgeRock site at [how to get involved](#).

Chapter 1

Solutions Overview

ForgeRock's Open Identity Gateway (OpenIG) provides the answer to three important challenges.


1.1 Extend SSO To Any Application

In today's enterprise, more than 30% of web applications are incompatible with web access management (WAM) software. That is, unlike OpenAM with OpenIG, most web access management products lack the agent to protect the web applications, or the application is a legacy solution that does not follow standard protocols for single sign-on. This limits the return on the enterprise WAM investment and constrains what types of web applications can be protected.

ForgeRock's Open Identity Gateway addresses this problem by extending access management to encompass *all* web applications. With OpenIG, OpenAM deployments can now be extended to be inclusive of those applications that do not integrate with policy agents alone. In addition, ForgeRock Open Identity Gateway interoperates, out-of-the-box, with all management solutions. Most importantly, your organization can on-board any web application without ever modifying or touching the target application again, significantly reducing the development and quality assurance required to protect web applications.

1.2 Federate Enabling Applications

The expertise and cost required to add SAML2 support to web applications is a problem for many businesses. Those businesses not moving to a standard for exposing their applications to their customers see increased cost and maintenance due to the complexity of one-off proprietary integrations. They may also see a loss of business to those customers requiring a Federation standard for authentication. Sometimes, deploying a full access management solution just to federate a few applications is too complex and costly, and building out their own solution by modifying their applications is just not possible.




OpenIG Federation Gateway allows businesses to add SAML2 support to their applications with little to no knowledge of the standard. In addition, there is no need to modify the application or install any plugin or agent on the application container.

1.3 Implement Standards Based Policy Enforcement

In order to provide SSO, OpenAM policy enforcement points, the policy agents, have traditionally functioned as plugins for each application. The policy agents however use OpenAM specific APIs and protocols. In contrast, you can roll out OpenIG as an independent policy enforcement point that is built on the SAML2 standard.

Implement Standards Based Policy Enforcement



OpenIG Federation Gateway can replace policy agents with a standards based SAML2 compliant enforcement point.

Chapter 2

How OpenIG Works

This chapter provides a detailed look at OpenIG's components and how they work together.

2.1 How It Works In A Nutshell

The underlying core of ForgeRock Open Identity Gateway is based on a [reverse proxy architecture](#). All HTTP traffic to each protected application is routed through OpenIG, enabling close inspection, transformation and filtering of each request. By inspecting the traffic, OpenIG is able to intercept requests that would normally require the user to authenticate, obtain the user's login credentials, and send the necessary HTTP request to the target application, thereby logging in the user without modifying or installing anything on the application. In its simplest form and basic configuration, OpenIG is a Java-based reverse proxy which runs as a web application. Enable the Form-Filter replay module and OpenIG automatically log users in when a timeout or authentication page is detected. Additionally, enable the SAML2 service and OpenIG becomes a SAML2 endpoint. In this mode of operation, OpenIG receives and verifies the SAML2 request and then logs the user directly into the target application.

2.2 How It Works In Detail

The following modules make up OpenIG.

2.2.1 OpenIG Core (Reverse Proxy)

OpenIG core is a standard Java EE servlet implementation of a reverse proxy. The main function of OpenIG core is to act as a reverse proxy to the target application. When deployed in its base configuration, OpenIG can be used as a pure reverse proxy. The power of the OpenIG core comes in its ability to search, transform, and process HTTP traffic to and from the target application. This enables OpenIG to recognize login pages, submit login forms, transform or filter content, and even function as a Federation endpoint for the application. All these features are possible without making any changes to the application's deployment container or the application itself.

2.2.2 Exchange

The [Exchange](#) is a wrapper around the HTTP request and response objects that pass through OpenIG. Every request or response being processed in OpenIG can be accessed or modified through the Exchange object. In addition, arbitrary data can be set in the Exchange to facilitate the passing of data and state between filters and handlers.

2.2.3 Dispatcher

The Dispatcher may be thought of as the internal router of OpenIG. Every request that comes into OpenIG is analyzed and forwarded on to the configured processing chain of filters and handlers. A request may be forwarded based on the target host, URL, URL parameters, headers, cookies, or any other component of the request.

2.2.4 Chain

A [Chain](#) is a combination of one or more Filters and a handler that process an incoming request from the Dispatcher. For example, the Dispatcher can process an incoming request with a URL parameter of action=login and forward the request to the Login Chain. The Login Chain executes a list of Filters and then calls a Handler. The Handler sends the request on to the target application or to another Chain for further processing.

2.2.5 Handlers

The final processing of every Chain ends in a call to a Handler. A Handler can simply call another Chain or it can send the request on to the target application. The following Handlers are shipped with OpenIG:

- [ClientHandler](#): Sends the final request to the target application.

- [StaticResponseHandler](#): Used to send a response, such as a redirect, to a client during request processing.
- [SequenceHandler](#): Links together multiple handlers or chains during request processing.

2.2.6 Filters

Filters are responsible for processing HTTP requests in OpenIG. Filters can be chained together to act on the input stream coming from the browser, or the output stream returned back from the target application. A filter can do something as simple as logging the input and output stream or something more complex, such as processing login pages, fetching user attributes, or transforming content. There are multiple Filters shipped with OpenIG that can be combined in chains to provide very extensible request and response processing features. Custom filters can also be written using the Java SPIs. The following is a list of Filters shipped with OpenIG:

- [CookieFilter](#): The default behavior of OpenIG is to accept and forward all cookies. Since this is not always the desirable behavior, the CookieFilter, when configured, allows you to suppress, manage, and relay cookies.
- [CaptureFilter](#): Captures the HTTP requests being processed by OpenIG. Capture can be used for audit purposes and may also be very useful when analyzing an application or troubleshooting a misbehaving OpenIG. Logs are written to a flat file on the OpenIG host.
- [HeaderFilter](#): The default behavior of OpenIG is to accept and forward all headers. The HeaderFilter can be configured to add additional headers or remove headers on both the HTTP request and the response. It can also be configured to parse and set header values in OpenIG context to allow filters access to the header attributes. This feature is used most commonly when OpenIG is integrated with OpenAM and being fronted by a policy agent.
- [AssignmentFilter](#): Sets values in the HTTP request and response.
- [StaticRequestFilter](#): Creates and sends HTTP GET and POST requests. The request can be formed using parameters from previous processing or statically configured values.
- [EntityExtractFilter](#): Searches for specific content with the body of the requests. For example, it can be used to extract hidden form parameters in a login page, which are needed in the login request.
- [ExceptionHandler](#): Sends users to configured URLs when errors or exceptions occur during request processing or user interactions.
- [SwitchFilter](#): Conditionally diverts the exchange to another handler.

- [HttpBasicAuthFilter](#): Performs HTTP basic authentication to the target application per [RFC 2617](#).
- [FileAttributesFilter](#): Looks up attributes in a flat file with the specified key. The attributes are added to the exchange to be used by subsequent filters or handlers.
- [SqlAttributesFilter](#): Executes an SQL prepared statement with configured parameters. The result is added to the exchange to be used by subsequent filters or handlers.

2.2.7 Configuration

The configuration of OpenIG was designed to be very modular and self-contained. Each module within OpenIG stores its configuration in JSON representation, which is stored in flat files. The features of OpenIG can be configured by directly manipulating the JSON flat files.

2.2.8 Heaplets

Every OpenIG module which has JSON configuration also has a Heaplet associated with it. Each module's Heaplet is responsible for reading the JSON configuration and creating that module's configuration in OpenIG's runtime heap. Each module can then read its configuration from the heap as well as make shared configuration information available to other modules.

2.2.9 Services

The Services model was designed to allow the integration of value added features, such as Federation, without requiring them to be tightly bound to the OpenIG implementation. SAML2 is the add on service shipping with OpenIG today.

2.2.10 Federation Service (Federation Gateway)

When the Federation Service is enabled, OpenIG acts as the Service Provider in a Circle of Trust with a SAML2 compliant Identity Provider. The Federation service supports both IDP and SP initiated SAML2 profiles. The Federation Gateway may be a Service Provider in the classic Federation use case where the IDP and SP are different companies or domains. The OpenIG SP can also be a standards based alternative to an OpenAM policy agent for internal enterprise applications that want to move away from installing agents with web and web application servers.

Chapter 3

Getting Started

This chapter provides instructions to get OpenIG up and running on Jetty, configured to execute the WordPress samples at our hosted WordPress deployment. This allows you to quickly see how OpenIG works and provides hands on experience with a few key features. If you are looking for more general installation and configuration instructions for your deployment or would like to use a container other than Jetty, start with the chapter on [Installing OpenIG](#).

3.1 Before You Begin

For the remainder of this chapter, \$HOME is the location of the home directory of the user running the web application container, where Jetty, OpenIG and the Samples are unpacked.

Make sure you have the correct Java environment installed. If necessary, download and install [Sun Java SE JDK 6 Update 21](#) or later. At this time these are the only versions of any JDK supported with OpenIG.

3.2 Installation Steps

1. [Download](#) and if necessary unpack the gateway .war file.
2. Download and unzip the [OpenIG sample configuration files](#).
3. Download and unzip [Jetty 7.1.6 Hightide](#).
4. Copy the OpenIG .war file as root.war to the webapps/ directory in Jetty.

When copied to this location along with the `root.xml` file below, Jetty automatically deploys OpenIG on startup.

```
$ cp $HOME/gateway-*.war $HOME/jetty-hightide-7.1.6/webapps/root.war
```

5. Copy the `WordPressLogin.json` gateway configuration file to `$HOME/.ForgeRock/OpenIG/config.json`. By default, OpenIG looks for `config.json` in the `$HOME/.ForgeRock/OpenIG/` directory.

```
$ cp $HOME/forgerock-sample-configs/WordPressLogin.json  
$HOME/.ForgeRock/OpenIG/config.json
```

6. Create `$HOME/jetty-hightide-7.1.6/contexts/root.xml` with the following content. This tells Jetty to deploy OpenIG in the root context:

```
<?xml version="1.0" encoding="ISO-8859-1"?>  
<!DOCTYPE Configure PUBLIC "-//Jetty//Configure//EN"  
  "http://www.eclipse.org/jetty/configure.dtd">  
<Configure class="org.eclipse.jetty.webapp.WebAppContext">  
  <Set name="contextPath"></Set>  
  <Set name="war"><SystemProperty name="jetty.home"  
    default="."/>webapps/root.war</Set>  
  <Set name="extractWAR">true</Set>  
  <Set name="copyWebDir">false</Set>  
  <Set name="defaultsDescriptor"><SystemProperty name="jetty.home"  
    default="."/>etc/webdefault.xml</Set>  
</Configure>
```

7. If you are managing multiple applications in one domain you must turn on domain cookies in Jetty. This can be done by adding the following property to `$JETTY_HOME/contexts/root.xml` or to `$JETTY_HOME/etc/jetty.xml`. Note that any changes to `jetty.xml` impact all web applications running in the container.

```
<Get name="sessionHandler">  
  <Get name="sessionManager">  
    <Set name="sessionDomain">.forgerock.com</Set>  
  </Get>  
</Get>
```

8. Start Jetty in the background:

```
$ $HOME/jetty-hightide-7.1.6/bin/jetty.sh start
```

Or start Jetty in the foreground:

```
$ java -jar $HOME/jetty-hightide-7.1.6/start.jar
```

3.3 Configuration Steps

So far you have deployed the gateway web application in the root context of Jetty on port 8080. Since the gateway is a reverse proxy you must make sure all traffic from your browser to <http://demo.forgerock.com:8080/wordpress/> goes through the gateway. The quickest way to do this is to add an entry to your `/etc/hosts` file on UNIX systems or `%SystemRoot%\system32\drivers\etc\hosts` on Windows. See [Wikipedia](#) for more information on host files. If you are running OpenIG and the browser on the same host, add this entry to the hosts file on the OpenIG host.

```
127.0.0.1 demo.forgerock.com
```

If you are running the browser and OpenIG on separate hosts, you must add the IP address of the host running OpenIG. For example, if OpenIG is running on a host with IP Address of 10.0.1.5:

```
10.0.1.5 demo.forgerock.com
```

Tip

Some browsers cache IP address resolutions, even after clearing all browsing data. It might be necessary to restart the browser if you are changing IP addresses of named hosts. The simplest way to make sure you have configured your DNS or host settings properly is to stop OpenIG and make sure you cannot reach the target application. If you can still reach it, you should double check your host settings. You should also make sure your name resolution is configured to check your host files before DNS. This configuration can be found in `/etc/nsswitch.conf` for most UNIX systems. Make sure `files` is listed before `dns`.

3.4 Try It Out

<http://demo.forgerock.com:8080/wordpress> should take you to the WordPress portal hosted at ForgeRock. When you click the login link, you are automatically logged in by OpenIG with the user name `george`.

What's happening behind the scenes?

When your browser goes to `http://demo.forgerock.com:8080/wordpress/`, it is actually connecting to OpenIG deployed on Jetty. OpenIG is configured to proxy all traffic it receives to the target application at `http://demo.forgerock.com:8080/wordpress/`. For this example, OpenIG is watching for the login page from WordPress. When it sees the login page, it creates and POSTs the login form for WordPress on behalf of the user. Since this is just a sample, the credentials for the user are hard coded in the gateway configuration file. In a real deployment, the credentials may come from a database, directory, flat file, HTTP headers, SAML assertion, or an OpenAM policy agent.

Chapter 4

Installing OpenIG

Before you install OpenIG, make sure you have the right version of Java, and install and configure the web container.

- Make sure you have the correct Java environment installed.

```
$ java -version
java version "1.6.0_24"
Java(TM) SE Runtime Environment (build 1.6.0_24-b07-334-9M3326)
Java HotSpot(TM) 64-Bit Server VM (build 19.1-b02-334, mixed mode)
```

- Refer to the chapter on [Configuring Deployment Containers](#) to ensure you have properly installed and configured your web application container and are running a supported version of Java.

For the remainder of this chapter \$HOME refers to the home directory of the user running the web container where you install OpenIG.

Procedure 4.1. Download, Install, & Configure

1. [Download](#) the OpenIG .war file.
2. Deploy the .war to the root context of the web application container.
3. Copy the [proxy and capture](#) template. You can jump straight to a configuration template that best matches your application, but it is highly recommended you start with the basic proxy template.

-
4. Copy the template to `$HOME/.ForgeRock/OpenIG/config.json`. Replace the `TARGETIP` tag in `config.json` with the IP address and port number of your target application.
 5. Modify DNS or host file settings so the name of your target application resolves to the IP address of OpenIG. Be sure to restart your browser after making this change.
 6. Start the web container where OpenIG is deployed and browse to the target application. OpenIG should now proxy all traffic to the application.
 7. To ensure the browser is going through OpenIG, stop the OpenIG web container, and verify you cannot reach the application. You can also look in the OpenIG capture log to verify the traffic is going through OpenIG. The location for this log is set in `config.json`, by default `/tmp/gateway.log`.

Chapter 5

Configuring Deployment Containers

This section provides installation and configuration tips. The following containers are supported.

- Tomcat 6.x and 7.x
- Jetty 6.x and 7.x
- Glassfish 2.x and 3.x
- JBoss 5.x
- Weblogic 10.x

5.1 Tomcat

Tomcat reference links:

- [Tomcat Documentation](#)
- [Configuring Tomcat for HTTPS](#)
- [Hardening Tomcat for production](#)

Download and install [Sun Java SE JDK 6 Update 21](#) or later. At this time these are the only versions of any JDK supported with OpenIG.

Download and install [Tomcat 6](#) or [Tomcat 7](#).

\$TOMCAT_HOME refers to the location where you have installed Tomcat.

Configure Tomcat to listen on the same port and protocol of the application you are protecting with OpenIG. If your application listens on both an HTTP and an HTTPS port you must configure Tomcat to do so as well. If you are installing to run the samples you do not have to modify the Tomcat port configuration since both the Tomcat default and the samples use port 8080.

To configure Tomcat to use a port other than 8080 modify the defaults in \$TOMCAT_HOME/conf/server.xml. The quickest way to do this is to search for the default value of 8080 and replace it with the new port number.

5.1.1 Tomcat & MySQL/JNDI Configuration

If OpenIG accesses an SQL database you must [configure Tomcat for JNDI](#). Once you have Tomcat configured you must add the jar for the database, and set up a JNDI data source and a reference to that data source in your web application. The following steps are for the MySQL Connector/J.

Download the [MySQL JDBC Driver Connector/J](#). Add it to the class path by copying it to \$TOMCAT_HOME/lib.

Add a JNDI data source for your MySQL server and database in \$TOMCAT_HOME/conf/context.xml.

```
<Resource name="jdbc/forgerock" auth="Container" type="javax.sql.DataSource"
maxActive="100" maxIdle="30" maxWait="10000"
username="admin" password="11111111" driverClassName="com.mysql.jdbc.Driver"
url="jdbc:mysql://localhost:3306/databasename"/>
```

Add a resource reference to the data source in \$TOMCAT_HOME/conf/web.xml.

```
<resource-ref>
  <description>MySQL Connection</description>
  <res-ref-name>jdbc/forgerock</res-ref-name>
  <res-type>javax.sql.DataSource</res-type>
  <res-auth>Container</res-auth>
</resource-ref>
```

5.1.2 Tomcat & Cookie Domains

If you use OpenIG for more than a single application and those applications are on different hosts, you must configure Tomcat to set domain cookies. To do this add the following to your \$TOMCAT_HOME/conf/Catalina/server/root.xml file.

```
<Context sessionCookieDomain=".forgerock.com" />
```


5.1.3 Tomcat & SSL

To get Tomcat up quickly on an SSL port you can add an entry similar to the following in `$TOMCAT_HOME/conf/server.xml`.

```
<Connector port="443" protocol="HTTP/1.1" SSLEnabled="true"
    maxThreads="150" scheme="https" secure="true"
    address="10.0.1.6"
    clientAuth="false" sslProtocol="TLS" keystorePass="11111111"/>
```

You must also have to make sure you have create a key store along with a self-signed certificate. You can do that with the following UNIX command:

```
$ keytool -genkey -alias tomcat -keyalg RSA
```

5.2 Jetty

Jetty reference links:

- [Jetty Documentation](#)
- [Running Jetty on port 80](#)
- [Deploying Web Applications](#)

Download and install [Sun Java SE JDK 6 Update 21](#) or later. At this time these are the only versions of any JDK supported with OpenIG.

Download and install [Jetty](#). The install directory is referred to as `$JETTY_HOME` from now on.

Configure Jetty to listen on the same port and protocol of the application you are protecting with OpenIG. If you are configuring OpenIG for the samples you do not have to modify the Jetty port configuration since the default for Jetty is 8080 and the samples use port 8080.

To configure Jetty to use a port other than 8080 you must modify the default `$JETTY_HOME/etc/jetty.xml`. The quickest way to do that is to search for the default value of 8080 and replace it with the new port number. The simplest way to get Jetty running on ports below 1024 is to run the Jetty process as root. This is only appropriate for sample purposes. Do not run Jetty as root in production.

If the application your are protecting is configured on an HTTPS port, you must configure Jetty for HTTPS as well.

If you are managing multiple applications in one domain you will need to turn on domain cookies in Jetty. This can be done by adding the following property to `$JETTY_HOME/contexts/root.xml` or to `$JETTY_HOME/etc/jetty.xml`. Note that any changes to `jetty.xml` impact all web applications running in the container.

```
<Get name="sessionHandler">
  <Get name="sessionManager">
    <Set name="sessionDomain">example.com<Set>
  <Get>
<Get>
```

If OpenIG accesses an SQL database you must [configure Jetty for JNDI](#) or use the [Jetty Hightide distribution](#), which comes pre-configured for JNDI. Once you have Jetty configured you must add the jar for the database, and set up a JNDI data source and a reference to that data source in your web application. The following steps are for the MySQL Connector/J.

Download the [MySQL JDBC Driver Connector/J](#) Add it to the Jetty class path by copying it to \$JETTY_HOME/lib/jndi.

Add a JNDI data source for your MySQL server and database in \$JETTY_HOME/etc/jetty.xml.

```
<New id="jdbc/forgerock" class="org.eclipse.jetty.plus.jndi.Resource">
  <Arg><Arg>
  <Arg>jdbc/forgerock<Arg>
  <Arg>
    <New class="com.mysql.jdbc.jdbc2.optional.MysqlConnectionPoolDataSource">
      <Set name="Url">jdbc:mysql://localhost:3306/databasename<Set>
      <Set name="User">mysqladmin<Set>
      <Set name="Password">11111111<Set>
    <New>
  <Arg>
<New>
```

Add a resource reference to the data source in \$JETTY_HOME/etc/webdefault.xml.

```
<resource-ref>
  <description>My DataSource Reference</description>
  <res-ref-name>jdbc/forgerock</res-ref-name>
  <res-type>javax.sql.DataSource</res-type>
  <res-auth>Container</res-auth>
</resource-ref>
```

5.3 JBoss

JBoss reference links:

- [JBoss 5 Documentation](#)
- [JBoss 6 Documentation](#)
- [Admin Console Quick Start Guide](#)

Download and install [Sun Java SE JDK 6 Update 21](#) or later. At this time these are the only versions of any JDK supported with OpenIG.

Download and install the community version of [JBoss 5.1](#). `$JBOSS_HOME` refers to the location you have installed JBoss. The assumption in this chapter is that you run the default server configuration.

Configure JBoss to listen on the same port and protocol of the application you are protecting with OpenIG. If your application listens on both an HTTP and an HTTPS port you must configure JBoss to do so as well. If you are installing to run the samples you do not have to modify the JBoss port configuration since both the JBoss default and the samples use port 8080.

To configure JBoss to use a port other than 8080 modify `server.xml` under your server configuration. For example, if you are using the default server configuration, `server.xml` is `$JBOSS_HOME/server/default/deploy/jbossweb.sar/server.xml`. The quickest way to do this is to search for the default value of 8080 and replace it with the new port number.

If OpenIG accesses an SQL database you must configure JBoss for JNDI. Once you have JBoss configured you must add the jar for the database, and set up a JNDI data source and a reference to that data source in your web application. The following steps are for the MySQL Connector/J.

Download the [MySQL JDBC Driver Connector/J](#) Add it to the JBoss class path by copying it to `$JBOSS_HOME/server/default/lib`.

Add a JNDI data source for your MySQL server and database in `$JBOSS_HOME/server/default/deploy/mysql-ds.xml`.

```
<datasources>
  <local-tx-datasource>
    <jndi-name>jdbc/forgerock</jndi-name>
    <connection-url>jdbc:mysql://localhost:3306/databasename</connection-url>
    <driver-class>com.mysql.jdbc.Driver</driver-class>
    <user-name>admin</user-name>
    <password>11111111</password>
    <min-pool-size>5</min-pool-size>
    <max-pool-size>20</max-pool-size>
    <idle-timeout-minutes>5</idle-timeout-minutes>
    <exception-sorter-class-name>
com.mysql.jdbc.integration.jboss.ExtendedMysqlExceptionSorter
    <exception-sorter-class-name>
com.mysql.jdbc.integration.jboss.MysqlValidConnectionChecker
    <valid-connection-checker-class-name>
com.mysql.jdbc.integration.jboss.MysqlValidConnectionChecker
    <valid-connection-checker-class-name>
  </local-tx-datasource>
</datasources>
```

Add a resource reference to the data source in `$JBOSS_HOME/server/default/deployers/jbossweb.deployer/web.xml`.

```
<resource-ref>
  <description>MySQL Connection</description>
  <res-ref-name>jdbc/forgerock</res-ref-name>
  <res-type>javax.sql.DataSource</res-type>
```

```
<res-auth>Container<res-auth>  
<resource-ref>
```

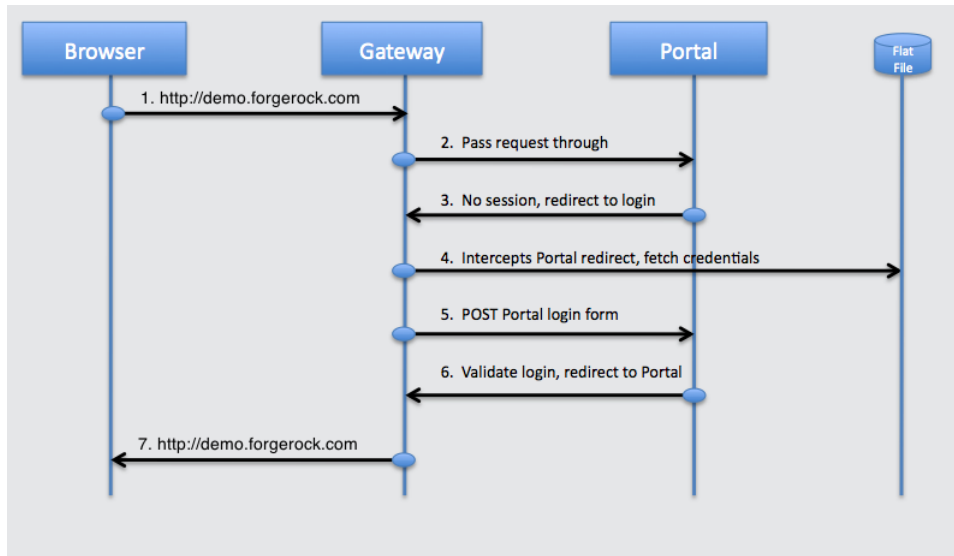
Chapter 6

Detailed Use Cases

This chapter diagrams and describes common use case request and response flows.

6.1 Portal Application Login

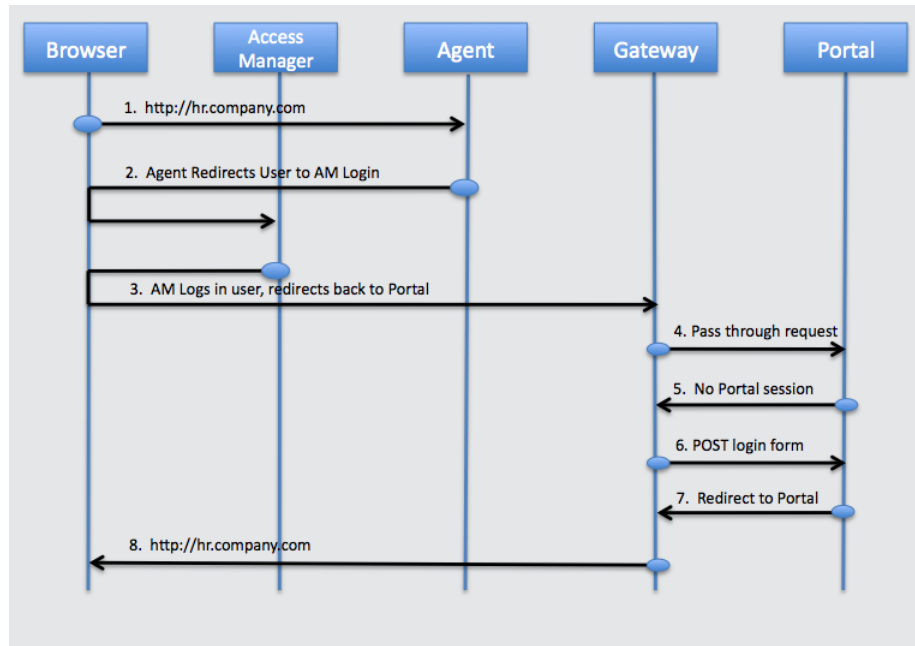
The figure below illustrates a sample flow with a description of each request from the browser to the back end application. This flow is taken from the [Tutorial With WordPress](#). Try the tutorial yourself to learn how OpenIG works. The Flat-File attribute store contains only one set of credentials. OpenIG makes the assumption this user is logging into the sample application. In a real deployment OpenIG would look up the user credentials using its own session, a SAML2 assertion, or a header from an OpenAM policy agent. Use cases that follow show examples of these types of deployments.



1. User accesses the Portal.
2. Gateway intercepts request, finds no login page and passes it through.
3. Portal finds no local session and redirects to its login page for authentication.
4. Gateway intercepts the redirect, finds a match for the login page, fetches the credentials from the flat file.
5. Gateway creates the login form and POSTs it to the Portal login page.
6. Portal validates the credentials and redirects to the user's home page.
7. Gateway passes the request through to the browser.

6.2 OpenAM Integration

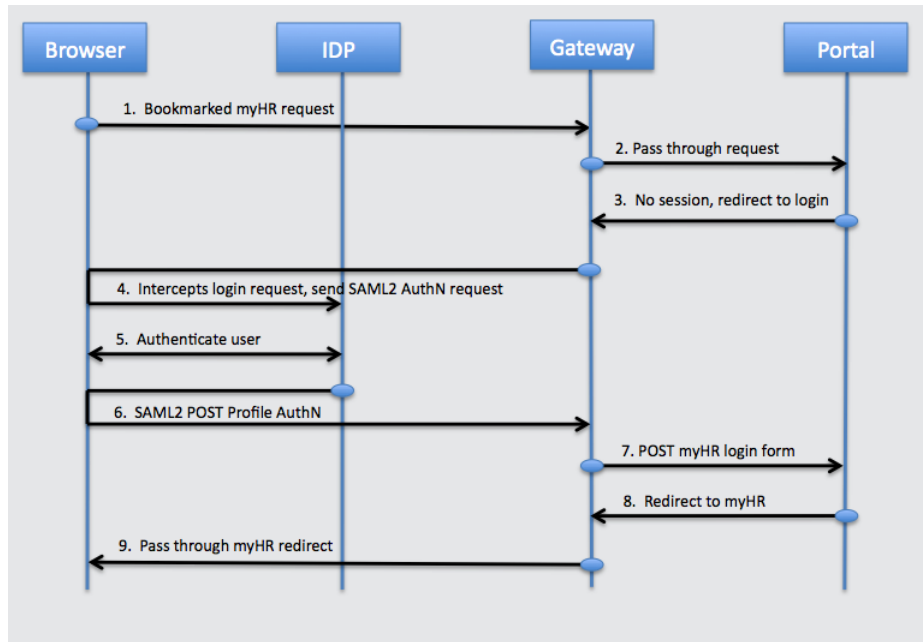
The figure below illustrates OpenIG integrated into an OpenAM deployment. In this deployment OpenIG is running in a container that is protected by an OpenAM policy agent. The agent is configured to forward a header, with the subject (user) of the single sign-on session, to OpenIG. OpenIG then uses the subject as the login credentials, or uses the subject as a reference to look up the login credentials in a database or directory. The HR application is integrated into the SSO deployment without an agent or any modification to the application or its deployment configuration.



1. User browses to the Portal.
2. OpenAM plugin intercepts the request, finds no valid OpenAM session, redirects the user to the OpenAM login pages.
3. OpenAM logs in the user and redirects back to the Portal.
4. OpenAM plugin finds a valid session, request goes through, OpenIG passes the request through to the Portal.
5. Portal finds no local session, redirects to the Portal login page.
6. OpenIG inspects the redirect, finds a match for the login page, creates the login form, and POSTs it to the Portal.
7. Portal validates the credentials and redirects to the Portal page.
8. Gateway passes the request through to the browser.

6.3 Federation Gateway SP Initiated SAML2 SSO

The figure below illustrates the OpenIG Federation Gateway providing SAML2 endpoint features acting as Service Provider in an SP initiated single sign-on configuration. In this sample, company myHR is an outsourced provider of HR services and has started seeing increased demand for SAML2 support in their core application. The companies to which they outsource are refusing proprietary means of authentication and demanding the widely accepted SAML2 standard. myHR is not in a position to modify their application to support SAML2 nor do they have the time or money to integrate and deploy all of OpenAM. With OpenIG Federation Gateway, myHR deploys OpenIG in front of their HR application, configures it as a SAML2 endpoint for SP initiated SSO, and configures it to log users into the HR application upon successful verification of the SAML2 assertion from their customers.

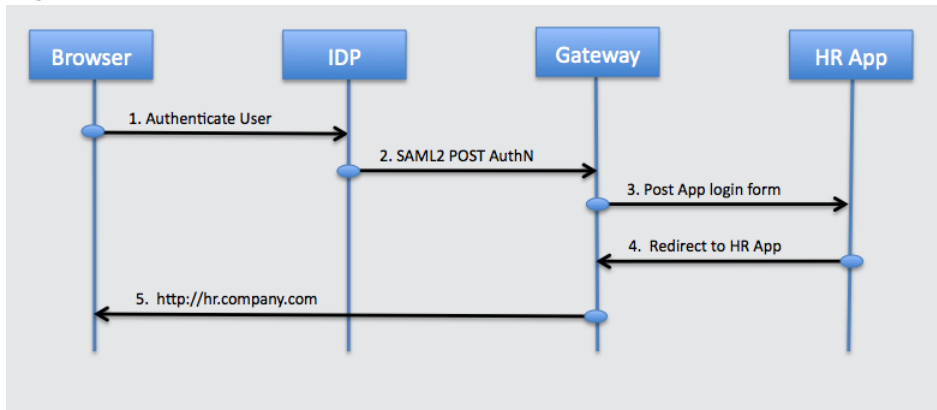


1. The user accesses the HR application through a bookmark in the browser.
2. Federation Gateway inspects the request, no match is found for the HR application's login page so the request goes through.
3. HR application finds no HR session, sends a redirect to its login page.
4. Federation Gateway intercepts the redirect, finds a match for the login page, issues an SP initiated SSO SAML2 request to Company.com's IDP.
5. The IDP at Company.com receives the SAML2 AuthN request and authenticates the user.

6. After authenticating the user the IDP sends a SAML2 POST to OpenIG.
7. Federation Gateway validates the assertion and makes the assertion attributes available to the OpenIG login chain.
8. OpenIG login chain gets the user credentials and POSTs the login form to the application.
9. The HR application verifies the credentials and redirects to its home page.

6.4 Federation Gateway IDP Initiated SAML2 SSO


The figure below illustrates the Federation Gateway providing SAML2 endpoint features acting as a Service Provider in a IDP initiated single sign-on configuration.



1. User clicks the HR link on the company portal and is redirected to the company IDP for authentication.
2. IDP sends an AuthN Response to the HR application.
3. Federation Gateway receives the POST, validates the assertion, and makes the attributes available to the OpenIG login chain.
4. OpenIG login chain retrieves the user credentials and POSTs the login form to the myHR application.
5. HR application validates the credentials and redirects to the main page of the application.

6.5 Standards Based SAML2 Agent

The figure below illustrates the Federation Gateway as a standards based replacement for OpenAM policy agents.



1. The user accesses an internal Payroll application.
2. Federation Gateway inspects the request, no match is found for the Payroll application's login page so the request goes through.
3. Payroll application finds no Payroll session, sends a redirect to its login page.
4. Federation Gateway intercepts the redirect, finds a match for the login page, issues an SP initiated SSO SAML2 request to OpenAM at Company.com.
5. OpenAM at Company.com receives the SAML2 AuthN and authenticates the user.
6. After authenticating the user OpenAM sends a SAML2 POST to OpenIG.
7. Federation Gateway validates the assertion and makes the assertion attributes available to the OpenIG login chain.
8. OpenIG login chain gets the user credentials and POSTs the login form to the application.
9. The Payroll application verifies the credentials and redirects to its home page.

Chapter 7

Tutorial With WordPress

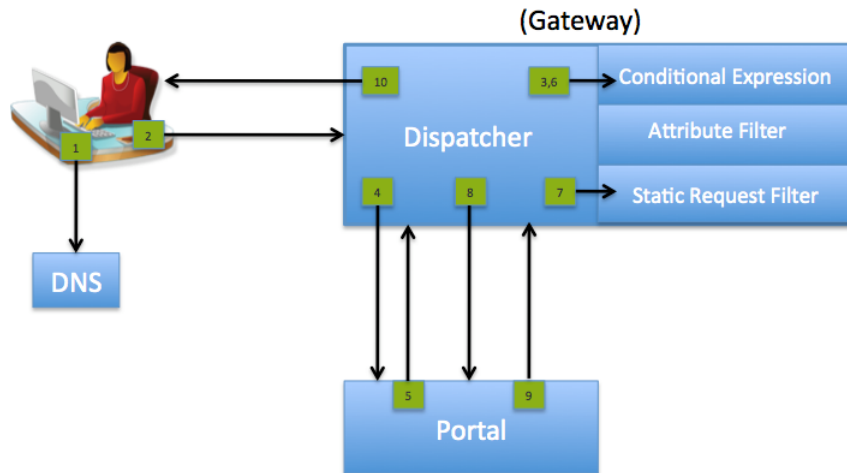
This tutorial gives you hands on experience with the basic features of OpenIG. You install the Gateway in your network and point it to an installation of the WordPress Portal hosted at <http://demo.forgerock.com:8080/wordpress>. The tutorial progresses through the following steps and explains what is happening by walking you through OpenIG configuration.

- Pure reverse proxy mode to the Portal
- Login with hard coded credentials
- Login with credentials from a flat file
- Login with credentials from MySQL

Before you begin the tutorial, read the chapter on *[How OpenIG Works](#)*.

7.1 Detailed Flow

The figure and the steps below detail the flow between the browser, OpenIG and the portal for a basic login sequence to the portal. The key concept to understand is all requests to the application are routed through OpenIG by modifying the DNS entry for the target application. OpenIG can then be configured to intercept specific pages, such as login pages, and create requests to log the user in.



1. User browses to the Portal, browser host makes a DNS request to get the IP address for `demo.forgerock.com`, DNS returns the IP address of OpenIG.
2. Browser sends the request to OpenIG.
3. OpenIG inspects the request, finds no login page match.
4. OpenIG forwards the original request to the Portal.
5. Portal finds no local application session and redirects to the Portal login page.
6. OpenIG intercepts the redirect, inspects the request, finds a login page match.
7. OpenIG creates the login form.
8. OpenIG POSTs the login form to the Portal.
9. Portal validates the login and redirects to the Portal end user page.
10. OpenIG receives the final redirect and passes it back to the Browser.

7.2 Setup Summary

To keep things simple, these instructions assume you are running OpenIG and the Browser on a single host. If they are on separate hosts, make sure the host or DNS changes impact the Browser host. This tutorial also assumes Jetty 7 for the container, but you can use any supported container.

- OpenIG IP address 10.0.1.1 (substitute your OpenIG's IP address)
- Jetty 7 listening on port 8080
- OpenIG deployed under the root context
- DNS or /etc/hosts pointing demo.forgerock.com to 10.0.1.1 (substitute your OpenIG's IP address)

7.3 Installing OpenIG

See [Getting Started](#) for instructions, and have your container listen on port 8080.

7.4 Configuring the Browser Host

With OpenIG running, configure the browser to go through OpenIG to get to the Portal. To do this we need to modify DNS or /etc/hosts so demo.forgerock.com points to the IP address of the host running OpenIG. If you are using /etc/hosts be sure to check that the naming services for your host check checking files before DNS. On UNIX systems this is configured in /etc/nsswitch.conf. Assuming OpenIG is running on 10.0.1.1, the /etc/hosts entry should be the following:

```
10.0.1.1 demo.forgerock.com
```

Before continuing please verify your name to IP address changes are working properly. The simplest way to do this is to ping demo.forgerock.com. You should see the following if you have properly configured your name to IP mapping:

```
$ ping demo.forgerock.com
ping demo.forgerock.com
PING demo.forgerock.com (10.0.1.1): 56 data bytes
64 bytes from 10.0.1.1: icmp_seq=0 ttl=64 time=0.053 ms
```

If the ping command returns the actual address of demo.forgerock.com, you need to recheck your settings.

7.5 Configuring OpenIG & Trying It Out

For the remainder of this document, \$HOME refers to the location where you unzipped your sample configuration files.

Download and unzip the [sample configuration files](#). If you installed through quick start you already have the configuration files downloaded and installed.

Copy `WordPressProxyOnly.json` to `$HOME/.ForgeRock/OpenIG/config.json`. By default, OpenIG looks for `$HOME/.ForgeRock/OpenIG/config.json` for its configuration. You must restart the OpenIG container when making any change to the configuration file.

```
$ cp $HOME/forgerock-sample-configs/WordPressProxyOnly.json
  $HOME/.ForgeRock/OpenIG/config.json
$ jetty.sh restart
```

To try out the first sample browse to <http://demo.forgerock.com:8080/wordpress>. You should see the WordPress Portal home page. If you click on the login link you should be prompted to login to the application as if you were accessing it directly. To verify you are actually going through OpenIG, stop the OpenIG container, refresh your browser and try to access the application again. If you still see the application, make sure your DNS or host files are configured to point to OpenIG instead of the Portal. You can login to the Portal with user name `george` and password `costanza`. The next section shows how to configure OpenIG to intercept the login page and automatically log you in when it sees that you have clicked the login link.

To see what is happening behind the scenes, take a look at `$HOME/.ForgeRock/OpenIG/config.json`. Look for the `HandlerServlet`. This is the servlet entry point to OpenIG. The `HandlerServlet` passes the request off to another handler which may be a chain of filters and handlers. In the pure proxy case there is no special logic to execute so it hands off to the `ClientHandler`. The job of the `ClientHandler` is to send the request on to the target. Since there are no chains called before the `ClientHandler`, the request passes through to the target untouched.

7.6 Login with Hard Coded Credentials

Now that OpenIG can proxy all traffic to and from the application, the next step is to configure OpenIG to intercept the login page, and POST the login form to the Portal.

```
$ cp $HOME/forgerock-sample-configs/WordPressLogin.json
  $HOME/.ForgeRock/OpenIG/config.json
$ jetty.sh restart
```

As you changed the OpenIG configuration file, you must restart the OpenIG container. After restarting you should be able to go to <http://demo.forgerock.com:8080/wordpress>, click on the login link, and be logged in as the user `george` without entering your credentials. OpenIG has intercepted the request for the login page, created the login form, and POSTed the request to the Portal.

To see what is happening behind the scenes, take a look at `$HOME/.ForgeRock/OpenIG/config.json`. This time you notice the `HandlerServlet` is calling the `DispatchHandler`. The `DispatchHandler` has a condition which checks for the presence of `wp-login.php` (the login page) in the URI path. If the condition is false (no login page), the `ClientHandler` is called sending the request on to the Portal. If the condition is true, OpenIG has found the login page, and calls the `LoginChain` for further processing. The `LoginChain` calls its filter, the `LoginRequestFilter`, which creates the login form, and then calls the `ClientHandler` to send the form to the Portal. If you look at the `LoginRequestFilter`, you notice it defines the method, URI, and form parameters for the request which is sent by the `ClientHandler`. In the form parameters you see the hard coded credentials, which are used to log you in as the user `george`.

7.7 Login with Credentials From a Flat File

Hard coding login credentials is great for a sample, but not realistic when it comes to a production deployment. In this next section you see how OpenIG can be configured to fetch the user's credentials from an external source, such as a directory or a database. For this sample, you use a flat file and the `FileAttributesFilter`. The key to look up the user in the flat file is hard coded, but you can understand how the key could come from something in the incoming request, such as an HTTP header, a session cookie, or even a SAML assertion.

For this sample you must modify the file attribute in the `FileAttributesFilter` object in `config.json` to reflect the location of the sample flat file. The location of the user file shipped with the samples is `forgerock-sample-configs/userfile`.

```
$ cp $HOME/forgerock-sample-configs/WordPressLoginFile.json  
$HOME/.ForgeRock/OpenIG/config.json  
$ jetty.sh restart
```

After restarting the container, check that you can go to <http://demo.forgerock.com:8080/wordpress> and be logged in as the user `george` without seeing the login page. To verify the login credentials are being picked up from the flat file, change `george`'s password, restart OpenIG, and try the login page again. You should get a login failed page. If you would like to log in as a different user, look for the `value` attribute in `config.json`, replace `george@seinfeld.com` with `kramer@seinfeld.com`. Both `george` and `kramer` have accounts on the Portal so you can log in with either.

To see what is happening behind the scenes, take a look at `$HOME/.ForgeRock/OpenIG/config.json`. For this sample, you added one additional filter to the `LoginChain` and made a slight modification to the `LoginRequestFilter`. Take a look at the `LoginChain` and you can see the `filters` attribute now has the value of `["FileAttributesFilter","LoginRequestFilter"]`. When the `LoginChain` is executed the `FileAttributesFilter` will be called prior to the `LoginRequestFilter`. The `FileAttributesFilter` sets the results of its lookup in the `Exchange` accessible

through the [Expressions](#) `${exchange.credentials.xxx}` where `xxx` is the name of any one of the values from the flat file. Now take a look at the form attribute in the `LoginRequestFilter`. The hard coded values for log and password have been replaced with `["${exchange.credentials.username}"]` and `["${exchange.credentials.password}"]`.

7.8 Login with Credentials From MySQL

This sample fetches the login credentials from MySQL. You must have access to an instance of MySQL in order to execute this sample. If you do not have access or just want to see how the `SqlAttributesFilter` works, just read the rest of this section.

For this sample you must configure your container for JNDI and MySQL. See the chapter on [Configuring Deployment Containers](#) for instructions on configuring containers for use with JNDI and MySQL. You also must create an entry for a user and create a query which returns a result with user name and password attributes. The query in the sample assumes there is a table named `users`, with attributes `username` and `password`. The parameter passed into the prepared statement is `email=george@seinfeld.com`.

```
$ cp $HOME/forgerock-sample-configs/WordPressLoginSQL.json  
$HOME/.ForgeRock/OpenIG/config.json  
$ jetty.sh restart
```

To see what is happening behind the scenes, take a look at `$HOME/.ForgeRock/OpenIG/config.json`. For this sample, you replaced the `FileAttributesFilter` with the `SqlAttributesFilter`. The only difference between the filters is how they retrieve the attribute value pairs to set in the [Exchange](#). Once the values are set in the Exchange, the `LoginRequestFilter` accesses them in the exact same way.

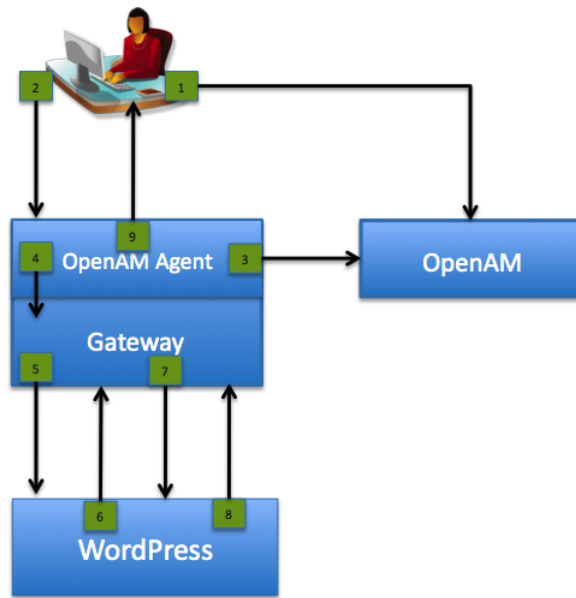
Chapter 8

Tutorial On OpenAM Password Capture & Replay

This tutorial walks you through an OpenAM integration with OpenAM's password capture and replay feature. This feature of OpenAM is typically used to integrate with Outlook Web Access or Sharepoint by capturing the password during OpenAM authentication, encrypting it, adding to the session, which is later decrypted and used for Basic Authentication to OWA or Sharepoint. This tutorial shows how you can configure OpenIG to use the user name and password from the OpenAM Authentication to log the user into WordPress. This is also how you would achieve OWA or Sharepoint integration.

8.1 Detailed Flow

The figure below illustrates the flow of requests for a user logging into OpenAM and then being logged into WordPress with the user name and password from the OpenAM login session.



1. User authenticates to OpenAM establishing a session, OpenAM is configured to capture the password, encrypt it, and then store it in the user's session.
2. User browses to a WordPress page requiring authentication.
3. OpenAM Agent intercepts the request, validates the user session, and populates HTTP headers with the user name and encrypted password from the session.
4. Agent passes the original request with populated headers to OpenIG.
5. OpenIG inspects the request, does not find a match for the WordPress login page, so forwards the request on.
6. WordPress finds no application session and redirects to its login page.
7. OpenIG intercepts, sees the login page, creates the login form with the user name and decrypted password from the headers inserted by the Agent.
8. WordPress validates the login form, returns an application session and the user's home page.
9. WordPress home page returned to the user.

8.2 Setup Summary

OpenAM is installed on <http://www.idp.com:8888/openam>. WordPress hosted on the ForgeRock demo site at <http://demo.forgerock.com:8080/wordpress>, populated with a sample user george with password costanza. OpenIG is deployed on an instance of Tomcat protected by an OpenAM J2EE agent. The J2EE agent is configured for cross domain SSO, and to add the user name and replay password headers to the HTTP requests.

8.3 Setup Details

This section assumes you are familiar with the components involved.

8.3.1 OpenAM server

Install and configure OpenAM on <http://www.idp.com:8888/openam> with the default configuration. You can install OpenAM with your configuration, but be sure to substitute in the tutorial accordingly.

Create a sample user with user name george and password constanza.

To test, startup OpenAM and make sure you can login with the user name george and password constanza.

8.3.2 Tomcat Agent Profile

Create the J2EE agent profile with the following settings:

- Server URL <http://www.idp.com:8888/openam>
- Agent URL <http://demo.forgerock.com:8080/agentapp>
- Under Global settings change the Agent filter mode from ALL to SSO_ONLY
- Under Application > Session Attributes Processing change the Session Attribute Fetch Mode from none to HTTP_Header
- Under Application > Session Attributes Processing > Session Attribute Mappings add UserToken=username and sunIdentityUserPassword=password
- Under SSO > Cross Domain SSO select the Enabled checkbox. If you have installed OpenAM in the forgerock.com domain you do not need to enable cross domain SSO.

8.3.3 Password Capture Configuration

- In the OpenAM console under Access Control > Realm > Authentication click "All Core Settings" and add `com.sun.identity.authentication.spi.ReplayPasswd` to the Authentication Post Processing Classes.
- Run **`java -classpath amserver.jar:opensso-sharedlib.jar com.sun.identity.common.DESGenKey`** to generate a shared key for the OpenAM Authentication plugin and OpenIG. As an example, if you have OpenAM installed under \$HOME on Tomcat you would run the command **`java -classpath $HOME/tomcat7/webapps/openam/WEB-INF/lib/amserver.jar:$HOME/tomcat7/webapps/openam/WEB-INF/lib/opensso-sharedlib.jar com.sun.identity.common.DESGenKey`**.
- In the OpenAM console under Configuration > Servers and Sites click on your server name, go to Advanced and add `com.sun.am.replaypasswd.key` with the value of the key generated in the previous step.

You must restart the OpenAM server after changing Advanced properties.

8.3.4 Tomcat Installation

Install on OpenIG host listening on port 8080.

/etc/hosts or equivalent with a mapping of `demo.forgerock.com` to `127.0.0.1`. This example assumes the browser and OpenIG are on the same host. If the browser is on a different host you must create a mapping so that all requests to `demo.forgerock.com` are resolved to the OpenIG host.

To test, startup Tomcat and make sure you can browse to `http://demo.forgerock.com:8080` and get the Tomcat home page. If you get a different page, then you have not modified your host file settings correctly. `demo.forgerock.com` should resolve to either localhost or the host OpenIG where is running.

8.3.5 Policy Agent installation

Install on the OpenIG host on Tomcat at `http://demo.forgerock.com:8080`

OpenSSO server URL: `http://www.idp.com:8888/openam`

Install agent filter in global `web.xml` `true`

Agent URL: `http://demo.forgerock.com:8080/agentapp`

To test, startup Tomcat and browse to the request headers example at <http://demo.forgerock.com:8080/examples/servlets/servlet/RequestHeaderExample>. You should be redirected to OpenAM for authentication. Once logged in as user george with password constanza you should be redirected back to the examples page and be able to see the user name and replay password headers. The user name header should be george, the password header should be encrypted.

8.3.6 Gateway Configuration

See [Installing OpenIG](#) for instructions, and have your container listen on port 8080. Install OpenIG in the root context by removing the existing root context directory `$HOME/tomcat/webapps/ROOT/` and then renaming the OpenIG .war to `$HOME/tomcat/webapps/ROOT.war`.

Download the [WordPressDecryptHeaderLogin.json](#) configuration file, and use it to replace `$HOME/.ForgeRock/OpenIG/config.json`.

Search for and replace DESKEY with the key generated when you enabled the password capture feature in OpenAM.

Restart OpenIG after making changes to `config.json`.

8.3.7 Trying It Out

Browse to WordPress. If you are not already logged into OpenAM you should be redirected to the OpenAM login page. You should login with user name george and password constanza. After login you should be redirected back to the WordPress portal homepage. Now click on the login link and OpenIG logs you in as George Costanza.

To see what is happening behind the scenes, take a look at `$HOME/.ForgeRock/OpenIG/config.json`. Look for the `HandlerServlet`. This is the servlet entry point to OpenIG. Notice it is calling the `DispatchHandler`. The `DispatchHandler` has a condition which checks for the presence of `wp-login.php` (the login page) in the URI path. If the condition is false (no login page), the `ClientHandler` is called sending the request on to the WordPress. If the condition is true, the Gateway has found the login page, and calls the `LoginChain` for further processing. The `LoginChain` calls its filters, the `CryptoFilter`, which looks for the replay password header and decrypts it, the `LoginRequestFilter`, which creates the login form, and then calls the `ClientHandler` to send the form to the WordPress. If you look at the `LoginRequestFilter`, you notice it defines the method, URI, and form parameters for the request which is sent by the `ClientHandler`. In the form

parameters you see the user name and password retrieved from the user name and replay password headers in the HTTP request. These headers are created by the OpenAM agent from the user's OpenAM session and the replay password header is decrypted by the CryptoFilter prior to the LoginRequestFilter being executed.

Chapter 9

Using the Federation Gateway

The Federation component of OpenIG is a standards based authentication service used by OpenIG to validate a user and retrieve key attributes of the user in order to log them into the target applications.

There are two ways the Federation Service can be invoked:

1. IDP initiated SSO, where the remote Identity Provider sends an unsolicited Authentication statement to OpenIG
2. SP initiated SSO where OpenIG calls the Federation Service to initiate Federated SSO with the Identity provider

In either case, the job of the Federation Service is to validate the user and pass the required attributes to OpenIG to log the user into target applications.

See the [Tutorial For the Federation Gateway](#) for hands on experience with the Federation Gateway.

9.1 Installation Overview

This section is a summary of the steps needed to prepare OpenIG to act as a Federation endpoint for your target application.

- Install the OpenIG Federation .war.
- Configure OpenIG to log users into the target application. Getting this to work before configuring Federation makes the process much simpler to troubleshoot if anything goes wrong.

- Add Federation configuration to the OpenIG configuration file.
- Add the assertion and subject mappings, optional redirect URI and optional logout URI to the Federation configuration.
- Export the Identity Provider MetaData from the remote Identity Provider or use the meta data from the OpenAM generated fedlet.
- Import OpenIG metadata to your Identity Provider.

9.2 Configuration File Overview

The Federation service is configured by modifying the OpenIG `config.json` and Federation specific XML files located in the home directory of the effective user running the web application container housing OpenIG. The home directory of this user is referred to as `$HOME_DIR` in this document. By default, the Federation service looks in `$HOME_DIR/.ForgeRock/OpenIG/` for `config.json` and `$HOME_DIR/.ForgeRock/SAML/` for the Federation XML configuration.

The following is a description of the files:

`$HOME_DIR/.ForgeRock/OpenIG/config.json`

This is the core configuration file for OpenIG. You must configure both this file and the XML files specific to the Federation Service. The reason there are two sets of configuration files is the Federation Service is based on the openFed library from the OpenAM open source project. The openFed library and the Federation configuration files are taken directly from the OpenFM release. In order to get the Federation Service configured you must tag swap the XML files. If you are familiar with the workflow in the OpenAM console you can generate a Fedlet and directly copy the configuration files into `$HOME_DIR/.ForgeRock/SAML/`.

`$HOME_DIR/.ForgeRock/SAML/FederationConfig.properties`

Advanced features of the openFed library. The defaults suffice in most deployments.

`$HOME_DIR/.ForgeRock/SAML/gateway.cot`

Circle of trust for OpenIG and the Identity Provider.

`$HOME_DIR/.ForgeRock/SAML/idp.xml`

This file is not included with the Federation Service XML. It must be generated by the Identity Provider and copied into the configuration directory.

`$HOME_DIR/.ForgeRock/SAML/idp-extended.xml`

Standard extensions for the Identity Provider.


```
$HOME_DIR/.ForgeRock/SAML/sp.xml  
$HOME_DIR/.ForgeRock/SAML/sp-extended.xml
```

These are the standard metadata and metadata extensions for the OpenIG Federation Service.

9.3 Configuring the Federation Service

The simplest way to configure the OpenIG Federation Gateway is to use the OpenAM task wizard to generate a Fedlet and then use the Fedlet configuration files in the Gateway configuration directory. If you use the Fedlet configuration files, simply unpack `Fedlet.war` and copy all the files listed above into `$HOME_DIR/.ForgeRock/SAML/`. You do not have to modify the files to do basic IDP and SP initiated SSO with OpenIG. Note that the sample Federation Gateway `config.json` templates uses `/saml` as the URI so your Fedlet endpoint should be specified as `protocol://host.domain:port/saml`.

If you do not use the Fedlet wizard, then follow the instructions for the unconfigured Fedlet and then copy the Fedlet configuration files to the `$HOME_DIR/.ForgeRock/SAML/` directory. You must also export the metadata from the IDP and copy it to `idp.xml` in the same directory.

9.4 Config.json Sample

The following sample configuration is from the tutorial on federating WordPress. The sample configuration receives a SAML assertion from OpenAM and then logs the user into WordPress using the user name and password from the assertion.

The following excerpt of the Federation Service JSON is from the OpenIG `config.json` file. All fields are mandatory except `logoutURI`. The `logoutURI` is only necessary if your configuration uses the single logout feature.

```
{  
  "name": "FederationServlet",  
  "type": "org.forgerock.openig.saml.FederationServlet",  
  "config": {  
    "assertionMapping": {  
      "userName": "uid",  
      "password": "userPassword"  
    },  
    "subjectMapping": "subjectName",  
    "sessionIndexMapping": "sessionIndex",  
    "redirectURI": "/login",  
    "logoutURI": "/logout",  
    "assertionConsumerEndpoint": "fedletapplication",  
    "SPinitiatedSSOEndpoint": "SPInitiatedSSO",  
    "singleLogoutEndpoint": "fedletSlo"  
  }  
}
```

name

Name of the Federation Service within the OpenIG runtime. This value should not be modified.

type

Classname of the Federation Servlet. This value should not be modified.

assertionMapping

The assertionMapping defines how to transform the attributes from the incoming assertion to attribute value pairs in the session. Each entry in the assertionMapping is of the form *attributeName:assertionName*. The *attributeName* represents the name of the attribute set in the session. The *assertionName* is used to fetch the value from the incoming assertion, which becomes the value in the session. The following statements correspond to the sample shown above.

If the incoming assertion contains the statement:

```
uid = jojo
```

```
userPassword = 123456789
```

Then the following values are set in the session:

```
userName = jojo
```

```
password = 123456789
```

Notice that you must also modify attributeMap in the \$HOME_DIR/.ForgeRock/SAML/sp-extended.xml to match the assertion mapping configured in the IDP metadata.

subjectMapping

The value contained in the assertion subject is set as the value of the attribute subjectName in the session.

sessionIndexMapping

The IDP's sessionIndex for the user is sent in the assertion. The value contained in the assertion is set as the value of the attribute sessionIndex in the session.

redirectURI

The redirectURI should be set to the page the Form-Filter recognizes as the login page for the target application. This is how OpenIG and the Federation service work together to provide single sign-on. When OpenIG detects the target application's login page, it redirects to the Federation Service. Once the Federation Service validates the SAML exchanges with the IDP and sets the required session attributes, it redirect back to the target application's login page. This allows the Form-Filter to finish the job of logging in the user by retrieving the necessary user data from the session and creating a login form for the application.

logoutURI

The logoutURI should be set to the URI which logs the user out of the target application. This attribute is only needed if your application uses the single log-out feature of the Identity Provider.

assertionConsumerEndpoint

The default value of fedletapplication is the same as the Fedlet. If you modify this attribute you must change the metadata to match.

SPinitiatedSSOEndpoint

The default value is SPInitiatedSSO. If you modify this attribute you must change the metadata to match.

singleLogoutEndpoint

The default value of fedletSLO is the same as the Fedlet. If you modify this attribute you must change the metadata to match.

9.5 Example Settings

Application myportal requires a form with userName and password for login. The userName for myportal is the mail attribute at the user's Identity Provider. The password for myportal is the mailPassword at the Identity Provider. The incoming SAML2 assertion sent by the Identity Provider contains the mail and mailPassword attributes. The Federation Service validates the incoming assertion, sets the userName and password in the HttpSession to the values of mail and mailPassword from the assertion, and redirects the user to /myportal/login. The LoginRequest Filter then retrieves the login values from the session and creates the form to log the user into myportal.

```
{
  "name": "FederationServlet",
  "type": "org.forgerock.openig.saml.FederationServlet",
  "config": {
    "assertionMapping": {
      "userName": "mail",
      "password": "mailPassword"
    }
  }
}
```

```
        "redirectURI": "/myportal/login",
        "logoutURI": "/myportal/logout"
    }
}

{
  "name": "LoginRequest",
  "type": "StaticRequestFilter",
  "config": {
    "method": "POST",
    "uri": "https://10.10.0.5/login",
    "form": {
      "userName": [${exchange.session.userName}],
      "password": [${exchange.session.password}],
    }
  }
}
```

9.6 Identity Provider Metadata

The Identity Provider metadata must be in `$HOME_DIR/.ForgeRock/SAML/idp.xml` directory. See the documentation for your Identity Provider for instructions on how to get the metadata.

To export Identity Provider metadata from OpenAM, run a command such as the following.

```
$ ssoadm
export-entity
-u amadmin
-f /tmp/pass
-y http://www.idp.com:8080/openam
-m /tmp/idp.xml
```

Chapter 10

Tutorial For the Federation Gateway

This tutorial expands on the [Tutorial With WordPress](#) by showing you how to configure OpenIG as a SAML2 Federation endpoint to log the user into WordPress. When you complete this tutorial you will have configured OpenAM to send a SAML2 assertion to OpenIG and configured OpenIG to validate the assertion and to use the mail and employeeNumber attribute values from the assertion to log the user into WordPress.

10.1 Before You Start

Follow the [Tutorial With WordPress](#), and make sure you can login to WordPress before continuing this tutorial.

10.2 Configuring OpenAM

For this sample and the test URIs, OpenAM is configured on `http://www.idp.com:8888/openam`. You can of course use any configuration, but be sure to substitute the proper values for the test URLs.

Use the OpenAM Console task wizard to create a Fedlet with the Name gateway and the Destination URL `http://demo.forgerock.com:8080/saml`. Configure the Attributes Mapping to contain the mail and employeeNumber attributes.

Create a user with mail attribute george and employeeNumber costanza.

10.3 Configuring OpenIG For Federation

Download the [Federation to WordPress Sample Config File](#) and copy it to `$HOME/.ForgeRock/OpenIG/config.json`, overwriting the `config.json` from the WordPress tutorial.

Assuming `Fedlet.zip` from the previous step is found in `$HOME/openam/myfedlets/` `Fedlet.zip` extract the Fedlet configuration files using the following commands:

```
$ cd $HOME/openam/myfedlets
$ unzip Fedlet.zip
$ jar xvf Fedlet.war conf
```

Create a `$HOME/.ForgeRock/SAML/` directory, and copy the Fedlet configuration to the new directory.

```
$ mkdir $HOME/.ForgeRock/SAML
$ cp conf/* $HOME/.ForgeRock/SAML
$ ls $HOME/.ForgeRock/SAML
FederationConfig.properties  idp-extended.xml      sp-extended.xml
fedlet.cot                   idp.xml                sp.xml
```

10.4 Trying It Out

For IDP initiated SSO [click this link](#)

For SP initiated SSO either browse to <http://demo.forgerock.com:8080/wordpress/> and click the login link, or [click this link](#)

To understand what is happening behind the scenes, take a look at `$HOME/.ForgeRock/OpenIG/config.json`. The `HandlerServlet` is no longer the entry point to OpenIG processing. If you look for the `servletObject` at the end of the file you see the `HandlerServlet` has been replaced by the `DispatchServlet`. The `DispatchServlet` provides an internal routing mechanism based on the incoming URI. The `DispatchServlet` looks at the incoming URI and dispatches to the `FederationServlet` when the SAML URI is found, otherwise it dispatches to the `HandlerServlet` as in the previous sample. This allows the `FederationServlet` to process the incoming assertion, set the attributes from the assertion in the `HTTPSession`, which then allows the normal chain of filter processing, through the `HandlerServlet`, to pick up the credentials and log the user into the target application.

Chapter 11

Configuration Templates

This chapter contains templates of common configurations. Start with one of our templates and then modify to suit your deployment. Read the summary of each template to find the right match for your application. If you are not sure about the characteristics of your application, start with the basic Application Capture template. This template allows you to setup basic proxying and capture the traffic of the login sequence in a flat file, which then allows you to analyze the application and subsequently choose the right template or add your own configuration.

Note

- All templates have the [CaptureFilter](#) enabled by default. Remove the capture filter from the outgoing chain before running the gateway in production. Capturing is typically used only for initial development or debugging and may rapidly fill up your available disk space if left enabled.
- Substitute the TARGETIP tag with the IP address of your application.
- Modify the LoginRequest filter to match the form required for login by your target application.

11.1 Proxy & Capture

Proxies all requests and captures them in a flat file. Use this template if you need to analyze the traffic for your application. Simply change the baseURI to be that of the target application, restart OpenIG, and login to the application. The entire sequence is logged to the flat file.

```
{
  "heap": {
    "objects": [
      {
        "name": "HandlerServlet",
        "comment": "Transforms Servlet request to handler request.",
        "type": "HandlerServlet",
        "config": {
          "handler": "DispatchHandler"
        }
      },
      {
        "name": "DispatchHandler",
        "type": "DispatchHandler",
        "config": {
          "bindings": [
            {
              "condition": "${exchange.request.uri.scheme == 'http'}",
              "handler": "OutgoingChain",
              "baseURI": "http://TARGETIP"
            },
            {
              "condition": "${exchange.request.uri.path == '/login'}",
              "handler": "LoginChain",
              "baseURI": "https://TARGETIP"
            },
            {
              "handler": "OutgoingChain",
              "baseURI": "https://TARGETIP"
            }
          ]
        }
      },
      {
        "name": "LoginChain",
        "type": "Chain",
        "config": {
          "filters": [],
          "handler": "OutgoingChain"
        }
      },
      {
        "name": "OutgoingChain",
        "type": "Chain",
        "config": {
          "filters": [
            "CaptureFilter"
          ],
          "handler": "ClientHandler"
        }
      },
      {
        "name": "CaptureFilter",
        "type": "CaptureFilter",
        "config": {
```



```
        "captureEntity": false,
        "file": "/tmp/gateway.log"
    },
    {
        "name": "ClientHandler",
        "comment": "Sends all requests to remote servers.",
        "type": "ClientHandler",
        "config": {}
    }
]
},
"servletObject": "HandlerServlet"
}
```

11.2 Simple Login Form

Logs the user into the target application with hard-coded user name and password. This template intercepts the login page request and replaces it with the login form.

```
{
  "heap": {
    "objects": [
      {
        "name": "HandlerServlet",
        "comment": "Entry point that transforms Servlet request to handler request.",
        "type": "HandlerServlet",
        "config": {
          "handler": "DispatchHandler",
          "baseURI": "http://TARGETIP"
        }
      },
      {
        "name": "DispatchHandler",
        "type": "DispatchHandler",
        "config": {
          "bindings": [
            {
              "condition": "${exchange.request.uri.path == '/login'}",
              "handler": "LoginChain",
            },
            {
              "handler": "OutgoingChain",
            },
          ]
        }
      },
      {
        "name": "LoginChain",
        "type": "Chain",
        "config": {
          "filters": ["LoginRequest"],
          "handler": "OutgoingChain"
        }
      },
      {
        "name": "LoginRequest",
```

```
        "type": "StaticRequestFilter",
        "config": {
            "method": "POST",
            "uri": "https://TARGETIP/login",
            "form": {
                "USER": ["myusername"],
                "PASSWORD": ["mypassword"],
            }
        },
    },
    {
        "name": "OutgoingChain",
        "type": "Chain",
        "config": {
            "filters": ["CaptureFilter"],
            "handler": "ClientHandler"
        }
    },
    {
        "name": "CaptureFilter",
        "type": "CaptureFilter",
        "config": {
            "captureEntity": false,
            "file": "/tmp/gateway.log",
        }
    },
    {
        "name": "ClientHandler",
        "comment": "Responsible for sending all requests to remote servers.",
        "type": "ClientHandler",
        "config": {
        }
    }
]
},
"servletObject": "HandlerServlet",
}
```

11.3 Login Form With Cookie From Login Page

For applications that expect a cookie from the login page to be sent in the login request form. This templates allows the login page request to go through to the target, intercepts the response, then creates the login form and adds the intercepted cookie to the POST.

```
{
  "heap": {
    "objects": [
      {
        "name": "HandlerServlet",
        "comment": "Entry point that transforms Servlet request to handler request.",
        "type": "HandlerServlet",
        "config": {
          "handler": "DispatchHandler",
          "baseURI": "http://TARGETIP"
        }
      },
    ],
  },
}
```

Login Form With Cookie From Login Page

```
"name": "DispatchHandler",
"type": "DispatchHandler",
"config": {
  "bindings": [
    {
      "condition": "${exchange.request.uri.path == '/eum/login'}",
      "handler": "LoginChain",
    },
    {
      "handler": "OutgoingChain",
    },
  ]
}
},
{
  "name": "LoginChain",
  "type": "Chain",
  "config": {
    "filters": ["SwitchFilter"],
    "handler": "OutgoingChain"
  }
},
{
  "name": "SwitchFilter",
  "type": "SwitchFilter",
  "config": {
    "onResponse": [
      {
        "handler": "LoginRequestHandler"
      }
    ]
  }
},
{
  "name": "LoginRequestHandler",
  "type": "Chain",
  "config": {
    "filters": ["LoginRequest"],
    "handler": "OutgoingChain"
  }
},
{
  "name": "LoginRequest",
  "type": "StaticRequestFilter",
  "config": {
    "method": "POST",
    "uri": "https://TARGETIP/login",
    "form": {
      "USER": ["myusername"],
      "PASSWORD": ["mypassword"],
    },
    "headers": {
      "cookie": ["${exchange.response.headers['Set-Cookie']}[0]"],
    }
  }
},
{
  "name": "OutgoingChain",
  "type": "Chain",
  "config": {
    "filters": ["CaptureFilter"],
    "handler": "ClientHandler"
  }
}
```

```
    }  
  },  
  {  
    "name": "CaptureFilter",  
    "type": "CaptureFilter",  
    "config": {  
      "captureEntity": false,  
      "file": "/tmp/gateway.log",  
    }  
  },  
  {  
    "name": "ClientHandler",  
    "comment": "Responsible for sending all requests to remote servers.",  
    "type": "ClientHandler",  
    "config": {  
    }  
  }  
]  
,  
"servletObject": "HandlerServlet",  
}
```

11.4 Login Form With Extract Filter & Cookie Filter

For applications that return the login page when the user tries to access a page without a valid session. This template shows how to use the ExtractFilter to find the login page on the response and use the CookieFilter to ensure the cookies from the application are replayed on each request. The sample application in this template is OpenAM. If you change the TARGETIP:PORT to be the IP address of OpenAM, the TARGETDN:PORT to be the fully qualified name and port of OpenAM and modify USERNAME and PASSWORD in the LoginRequest you automatically log USERNAME into OpenAM.

Note

Without the CookieFilter in the OutgoingChain the cookie set in the login page response would not get set in the browser since that request is intercepted before it gets to the browser. The simplest way to deal with this situation is to let OpenIG manage all the cookies by enabling the CookieFilter. The side effect of OpenIG managing cookies is none of the cookies are sent to the browser, but are managed locally by OpenIG.

```
{  
  "heap": {  
    "objects": [  
      {  
        "name": "HandlerServlet",  
        "comment": "Entry point that transforms Servlet request to handler request.",  
        "type": "HandlerServlet",  
        "config": {  
          "handler": "FindLoginPageChain",  
          "baseURI": "http://TARGETIP:PORT"  
        }  
      }  
    ]  
  }  
}
```

```
    },
    {
        "name": "FindLoginPageChain",
        "type": "Chain",
        "config": {
            "filters": ["IsLoginPage", "FindLoginPage"],
            "handler": "OutgoingChain"
        }
    },
    {
        "name": "FindLoginPage",
        "type": "EntityExtractFilter",
        "config": {
            "messageType": "response",
            "target": "${exchange.isLoginPage}",
            "bindings": [
                {
                    "key": "found",
                    "pattern": "OpenAM\\s\\(Login\\)",
                    "template": "true"
                }
            ]
        }
    },
    {
        "name": "IsLoginPage",
        "type": "SwitchFilter",
        "config": {
            "onResponse": [
                {
                    "condition": "${exchange.isLoginPage.found == 'true'}",
                    "handler": "LoginChain"
                }
            ]
        }
    },
    {
        "name": "LoginChain",
        "type": "Chain",
        "config": {
            "filters": ["LoginRequest"],
            "handler": "OutgoingChain"
        }
    },
    {
        "name": "LoginRequest",
        "type": "StaticRequestFilter",
        "config": {
            "method": "POST",
            "uri": "http://TARGETIP:PORT/openam/UI/Login"
            "form": {
                "IDToken0": [""]
                "IDToken1": ["USERNAME"]
                "IDToken2": ["PASSWORD"]
                "IDButton": ["Log+In"]
                "encoded": ["false"]
            },
            "headers": {
                "host": ["TARGETFQDN:PORT"]
            }
        }
    }
}
```

Login Which Requires a Hidden Value From the Login Page

```
    },
    {
        "name": "OutgoingChain",
        "type": "Chain",
        "config": {
            "filters": ["CookieFilter", "CaptureFilter"],
            "handler": "ClientHandler"
        }
    },
    {
        "name": "CookieFilter",
        "type": "CookieFilter",
        "config": {
        },
    },
    {
        "name": "CaptureFilter",
        "type": "CaptureFilter",
        "config": {
            "captureEntity": true,
            "file": "/tmp/gateway.log",
        }
    },
    {
        "name": "LogSink",
        "comment": "Default sink for logging information.",
        "type": "ConsoleLogSink",
        "config": {
            "level": "DEBUG",
        }
    }
}

{
    "name": "ClientHandler",
    "comment": "Responsible for sending all requests to remote servers.",
    "type": "ClientHandler",
    "config": {
    }
}
]
},
"servletObject": "HandlerServlet",
}
```

11.5 Login Which Requires a Hidden Value From the Login Page

Extracts a hidden value from the login page and includes it in the login form POSTed to the target application.

```
{
  "heap": {
    "objects": [
      {
        "name": "HandlerServlet",
        "comment": "Entry point that transforms Servlet request to handler request.",
        "type": "HandlerServlet",
        "config": {
          "handler": "DispatchHandler",
        }
      }
    ]
  }
}
```

Login Which Requires a Hidden Value From the Login Page

```

        "baseURI": "http://TARGETIP"
    },
    {
        "name": "DispatchHandler",
        "type": "DispatchHandler",
        "config": {
            "bindings": [
                {
                    "condition": "${exchange.request.uri.path == '/login'}",
                    "handler": "LoginChain",
                },
                {
                    "handler": "OutgoingChain",
                },
            ]
        }
    },
    {
        "name": "LoginChain",
        "type": "Chain",
        "config": {
            "filters": ["HiddenValueExtract", "LoginRequest"],
            "handler": "OutgoingChain"
        }
    },
    {
        "name": "HiddenValueExtract",
        "type": "EntityExtractFilter",
        "config": {
            "messageType": "response",
            "target": "${exchange.hiddenValue}",
            "bindings": [
                {
                    "key": "value",
                    "pattern": "wpLoginToken\\|s.*value=\\(.*\\)\"",
                    "template": "$1"
                }
            ]
        }
    },
    {
        "name": "LoginRequest",
        "type": "StaticRequestFilter",
        "config": {
            "method": "POST",
            "uri": "https://TARGETIP/login",
            "form": {
                "USER": ["myusername"],
                "PASSWORD": ["mypassword"],
                "hiddenValue": ["${exchange.hiddenValue.value}"],
            }
        }
    },
    {
        "name": "OutgoingChain",
        "type": "Chain",
        "config": {
            "filters": ["CaptureFilter"],
            "handler": "ClientHandler"
        }
    },

```

```
{
  {
    "name": "CaptureFilter",
    "type": "CaptureFilter",
    "config": {
      "captureEntity": false,
      "file": "/tmp/gateway.log",
    }
  },
  {
    "name": "ClientHandler",
    "comment": "Responsible for sending all requests to remote servers.",
    "type": "ClientHandler",
    "config": {
    }
  }
]
},
"servletObject": "HandlerServlet",
}
```

11.6 HTTP & HTTPS Application

Proxies traffic to an application listening on ports 80 and 443. The assumption is the application uses HTTPS for authentication and HTTP for the general application features. Assuming the login will all take place on port 443, you will need to add the login filters and handlers to the LoginChain. To get started quickly, modify the baseURI to be the IPAddress of your target application. This should allow you to proxy all traffic to the application. Then add the logic for the LoginChain using the flow from one of the login templates.

```
{
  "heap": {
    "objects": [
      {
        "name": "HandlerServlet",
        "comment": "Entry point that transforms Servlet request to handler request.",
        "type": "HandlerServlet",
        "config": {
          "handler": "DispatchHandler",
        }
      },
      {
        "name": "DispatchHandler",
        "type": "DispatchHandler",
        "config": {
          "bindings": [
            {
              "condition": "${exchange.request.uri.scheme == 'http'}",
              "handler": "OutgoingChain",
              "baseURI": "http://TARGETIP"
            },
            {
              "condition": "${exchange.request.uri.path == '/login'}",
              "handler": "LoginChain",
              "baseURI": "https://TARGETIP"
            }
          ]
        }
      }
    ]
  }
}
```



```

        "handler": "OutgoingChain",
        "baseURI": "https://TARGETIP"
    },
    ],
    },
    {
        "name": "LoginChain",
        "type": "Chain",
        "config": {
            "filters": [],
            "handler": "OutgoingChain"
        }
    },
    {
        "name": "OutgoingChain",
        "type": "Chain",
        "config": {
            "filters": ["CaptureFilter"],
            "handler": "ClientHandler"
        }
    },
    {
        "name": "CaptureFilter",
        "type": "CaptureFilter",
        "config": {
            "captureEntity": false,
            "file": "/tmp/gateway.log",
        }
    },
    {
        "name": "ClientHandler",
        "comment": "Responsible for sending all requests to remote servers.",
        "type": "ClientHandler",
        "config": {
        }
    }
    ],
    "servletObject": "HandlerServlet",
}

```

11.7 Multiple Applications

Shows how a single OpenIG configuration can proxy to multiple applications on different IPs and ports. This template is setup to proxy to both the WordPress and MediaWiki hosted samples. OpenIG must be setup to listen on both ports 8080 and 8181. The DispatcherHandler is used to route each incoming request to either the WordPress chain or the MediaWiki chain.

```

{
  "heap": {
    "objects": [
      {
        "name": "HandlerServlet",
        "type": "HandlerServlet",
        "config": {

```

```

        "handler": "DispatchHandler"
    }
},
{
    "name": "DispatchHandler",
    "type": "DispatchHandler",
    "config": {
        "bindings": [
            {
                "condition": "${exchange.request.uri.host == 'demo.forgerock.com'}",
                "handler": "WordPressLoginChain",
                "baseURI": "http://109.73.67.52:8080",
            },
            {
                "condition": "${exchange.request.uri.host == 'demo.forgerock.com'}",
                "handler": "MediaWikiLoginChain",
                "baseURI": "http://109.73.67.52:8181"
            }
        ]
    }
},
{
    "name": "WordPressLoginChain",
    "type": "Chain",
    "config": {
        "filters": [],
        "handler": "OutgoingChain"
    }
},
{
    "name": "MediaWikiLoginChain",
    "type": "Chain",
    "config": {
        "filters": [],
        "handler": "OutgoingChain"
    }
},
{
    "name": "OutgoingChain",
    "type": "Chain",
    "config": {
        "filters": ["CaptureFilter"],
        "handler": "ClientHandler"
    }
},
{
    "name": "CaptureFilter",
    "type": "CaptureFilter",
    "config": {
        "captureEntity": false,
        "file": "/tmp/gateway.log"
    }
},
{
    "name": "ClientHandler",
    "comment": "Responsible for sending all requests to remote servers.",
    "type": "ClientHandler",
    "config": {
    }
}
],
},

```

```
"servletObject": "HandlerServlet",  
}
```

11.8 OpenAM Integration With Headers

Logs the user into the target application using the headers passed down from an OpenAM policy agent. This template assumes the user name and password are passed down by the OpenAM policy agent as headers. If the header passed in contains only a user name or subject and requires a lookup to an external data source, you must add an attribute filter to the chain to retrieve the credentials.

```
{  
  "heap": {  
    "objects": [  
      {  
        "name": "HandlerServlet",  
        "comment": "Entry point that transforms Servlet request to handler request.",  
        "type": "HandlerServlet",  
        "config": {  
          "handler": "DispatchHandler",  
          "baseURI": "http://TARGETIP"  
        }  
      },  
      {  
        "name": "DispatchHandler",  
        "type": "DispatchHandler",  
        "config": {  
          "bindings": [  
            {  
              "condition": "${exchange.request.uri.path == '/login'}",  
              "handler": "LoginChain",  
            },  
            {  
              "handler": "OutgoingChain",  
            }  
          ]  
        }  
      },  
      {  
        "name": "LoginChain",  
        "type": "Chain",  
        "config": {  
          "filters": ["LoginRequest"],  
          "handler": "OutgoingChain"  
        }  
      },  
      {  
        "name": "LoginRequest",  
        "type": "StaticRequestFilter",  
        "config": {  
          "method": "POST",  
          "uri": "https://TARGETIP/login",  
          "form": {  
            "USER": ["${exchange.request.headers['username']}[0]}",  
            "PASSWORD": ["${exchange.request.headers['password']}[0]"}  
          }  
        }  
      }  
    ]  
  }  
}
```

```
    },
    {
      "name": "OutgoingChain",
      "type": "Chain",
      "config": {
        "filters": ["CaptureFilter"],
        "handler": "ClientHandler"
      }
    },
    {
      "name": "CaptureFilter",
      "type": "CaptureFilter",
      "config": {
        "captureEntity": false,
        "file": "/tmp/gateway.log",
      }
    },
    {
      "name": "ClientHandler",
      "comment": "Responsible for sending all requests to remote servers.",
      "type": "ClientHandler",
      "config": {
      }
    }
  ]
},
"servletObject": "HandlerServlet",
}
```

11.9 Microsoft Online Outlook Web Access

A sample template used to log a user into Microsoft Online Outlook Web Access. This template shows how you would use the Gateway and the OpenAM password capture feature to integrate with OWA. You can follow the [Tutorial On Password Capture & Replay](#) tutorial and substitute this template.

```
{
  "heap": {
    "objects": [
      {
        "name": "LogSink",
        "comment": "Default sink for logging information.",
        "type": "ConsoleLogSink",
        "config": {
          "level": "DEBUG",
        }
      },
      {
        "name": "HandlerServlet",
        "type": "HandlerServlet",
        "config": {
          "handler": "DispatchHandler",
          "baseURI": "https://65.55.171.158"
        }
      },
      {
        "name": "DispatchHandler",
        "type": "DispatchHandler",
      }
    ]
  }
}
```

```
"config": {
  "bindings": [
    {
      "condition": "${exchange.request.uri.path == '/owa/auth/logon.aspx'}",
      "handler": "LoginChain",
    },
    {
      "handler": "OutgoingChain",
    }
  ]
},
{
  "name": "LoginChain",
  "type": "Chain",
  "config": {
    "filters": ["CryptoHeaderFilter", "LoginRequest"],
    "handler": "OutgoingChain"
  }
},
{
  "name": "CryptoHeaderFilter",
  "type": "CryptoHeaderFilter",
  "config": {
    "messageType": "REQUEST",
    "operation": "DECRYPT",
    "algorithm": "DES/ECB/NoPadding",
    "key": "DESKEY",
    "keyType": "DES",
    "charSet": "utf-8",
    "headers": ["password"],
  },
},
{
  "name": "LoginRequest",
  "type": "StaticRequestFilter",
  "config": {
    "method": "POST",
    "uri": "https://65.55.171.158/owa/auth/owaauth.dll",
    "headers": {
      "Host": ["red001.mail.microsoftonline.com"],
      "Content-Type": ["Content-Type:application/x-www-form-urlencoded"],
    }
    "form": {
      "destination": ["https://red001.mail.microsoftonline.com/owa/"],
      "forcedownlevel": ["0"],
      "trusted": ["0"],
      "username": ["${exchange.request.headers['username']}[0]"},
      "password": ["${exchange.request.headers['password']}[0]"},
      "isUtf8": ["1"],
    }
  }
},
{
  "name": "OutgoingChain",
  "type": "Chain",
  "config": {
    "filters": ["HeaderFilter", "CaptureFilter"],
    "handler": "ClientHandler"
  }
},
{
```

```
        "name": "HeaderFilter",
        "type": "HeaderFilter",
        "config": {
            "messageType": "REQUEST",
            "remove": ["password", "username"],
        }
    },
    {
        "name": "CaptureFilter",
        "type": "CaptureFilter",
        "config": {
            "captureEntity": false,
            "file": "/tmp/gateway.log",
        }
    },
    {
        "name": "ClientHandler",
        "type": "ClientHandler",
        "config": {
        }
    }
],
"servletObject": "HandlerServlet",
}
```

Chapter 12

Customizing OpenIG

OpenIG includes a complete, stable [application programming interface](#), designed to allow you to customize OpenIG as required. Customizing OpenIG can be necessary if you need to perform complex server interactions or intensive data transformations, and stock handlers, filters and [expressions](#) provided by the gateway are too basic to meet the task.

12.1 Key Extension Points

The two primary extension points are the interfaces: [Filter](#) (for processing a request and/or response en route) and [Handler](#) (for generating responses from requests). These interfaces are similar to the Java Enterprise Edition `Filter` and `Servlet` interfaces, with some differences in the semantics of messages. While you can simply implement these interfaces, there are also included convenience classes: [GenericFilter](#) and [GenericHandler](#) that you can use if you intend to make your extensions configurable through the OpenIG configuration resource.

OpenIG also allows standard Java Enterprise Edition servlets and servlet filters to be embedded in the gateway, with initialization parameters supplied through the OpenIG configuration resource.

12.2 Implementing a Filter

The Filter interface exposes a [filter](#) method, which takes an [Exchange](#) object and the [Chain](#) of remaining filters and handler to dispatch to. Initially, [exchange.request](#) contains the request to be filtered. To pass the request to the next filter or handler in the chain, the filter calls [chain.handle\(exchange\)](#). After this call, [exchange.response](#) contains the response that can be filtered.

A filter might elect not to pass the request to the next filter or handler, and instead handle the request itself. It can achieve this by merely avoiding a call to `chain.handle(exchange)` and creating its own response object in the exchange. The filter is also at liberty to replace a response with another of its own. A filter can exist in more than one chain, therefore should make no assumptions or correlations using the chain it is supplied. The only valid use of a chain by a filter is to call its `handle` method to dispatch the exchange to the rest of the chain.

Note

If an existing response exists in the exchange object and the filter intends to replace it with its own, it must first check to see if the existing response has an [entity](#), and if it does, must call its [close](#) method in order to signal that the processing of the response from a remote server is complete.

12.3 Implementing a Handler

The Handler interface exposes a [handle](#) method, which takes an [Exchange](#) object. It processes the request in [exchange.request](#) and produces a response in [exchange.response](#). A handler can elect to dispatch the exchange to another handler or chain.

Note

If an existing response exists in the exchange object and the filter intends to replace it with its own, it must first check to see if the existing response has an [entity](#), and if it does, must call its [close](#) method in order to signal that the processing of the response from a remote server is complete.

12.4 Heap Object Configuration

Objects are added to the heap and supplied with configuration artifacts at initialization time. To be integrated with the configuration, a class must have an accompanying implementation of the [Heaplet](#) interface. The easiest and most common way of exposing the heaplet is to extend the [NestedHeaplet](#) class in a nested class in the class you want to create and initialize and implementing its [create](#) method.

Within the create method, you can access the object's configuration through the [config](#) field.

12.5 Including Servlets and Servlet Filters

A servlet or servlet filter can be added as a heap object by including it in the classpath, and identifying it in the OpenIG heap configuration as an [javax.servlet.http.HttpServlet](#) or [javax.servlet.Filter](#) type, providing any appropriate initialization parameters. A servlet can be dispatched directly by the [gateway servlet](#), and servlets and filters can be flexibly dispatched to through a [DispatchServlet](#).

12.6 Sample Filter

The following sample filter sets an arbitrary header in the incoming request and outgoing response.

```
package com.example.filter;

// Java Standard Edition
import java.io.IOException;

// OpenIG Core Library
import org.forgerock.openig.filter.Chain;
import org.forgerock.openig.filter.GenericFilter;
import org.forgerock.openig.handler.HandlerException;
import org.forgerock.openig.heap.HeapException;
import org.forgerock.openig.heap.NestedHeaplet;
import org.forgerock.openig.http.Exchange;
import org.forgerock.openig.model.ModelException;

public class HelloWorldFilter extends GenericFilter {

    public String name;
    public String value;

    @Override
    public void filter(Exchange exchange, Chain chain)
        throws HandlerException, IOException {
        exchange.request.headers.put(name, value); // set header in request
        chain.handle(exchange); // pass to remaining filters & handler in chain
        exchange.response.headers.put(name, value); // set header in response
    }

    public static class Heaplet extends NestedHeaplet {
```

```
@Override
public Object create() throws HeapException, ModelException {
    HelloWorldFilter filter = new HelloWorldFilter();
    filter.name = config.get("name").required().asString(); // required
    filter.value = config.get("value").required().asString(); // req'd
    return filter;
}
}
```

The corresponding heap object configuration is as follows.

```
{
  "name": "HelloWorldFilter",
  "type": "com.example.filter.HelloWorldFilter",
  "config": {
    "name": "X-Hello",
    "value": "World"
  }
}
```

Chapter 13

Troubleshooting

This chapter covers common problems and their solutions.

13.1 Object not found in heap

```
HTTP ERROR 500
org.forgerock.openig.model.NodeException:
  ['file:/Users/george/.ForgeRock/OpenIG/config.json'].heap.objects[2]
  .config.filterObjects[0]: object not found in heap
at org.forgerock.openig.heaplet.HeapUtil.getRequiredObject(HeapUtil.java:54)
at org.forgerock.openig.filter.Chain$Heaplet.create(Chain.java:49)
```

You have a filter specified in the filterObjects list in the Chain object that is not in config.json. Make sure you have added an entry for the Filter and have correctly spelled its name in the filterObjects list.

13.2 Could not find local configuration file

```
HTTP ERROR 500
Problem accessing /. Reason:

could not find local configuration file at
/var/root/.ForgeRock/OpenIG/config.json or bootstrap file at
/var/root/.ForgeRock/OpenIG/_private_var_folders_...webinf.json
```

OpenIG could not find its configuration file in the home directory of the user running the container where OpenIG is deployed, in this case `/var/root/.ForgeRock/OpenIG/config.json`. As you can see from the error message, OpenIG looks in a secondary location determined at runtime by the container. The secondary location is only necessary if you are running more than one OpenIG on a single host and your instances do not share a single configuration.

13.3 Unexpected character (x) at position 1103

```
HTTP ERROR 500
Problem accessing /. Reason:

Unexpected character (x) at position 1103
```

This error usually means a missing double quote or a missing bracket in the configuration file. Use a JSON editor or JSON validation tool such as [JSONLint](#) to make sure your JSON is valid.

13.4 The values in the flat file are incorrect

Ensure the flat file is readable by the user running the container for OpenIG. Values are all characters, including space and tabs, between the separator, so make sure the values are not padded with spaces.

13.5 Problem accessing URL

```
HTTP ERROR 500

Problem accessing /myURL . Reason:

java.lang.String cannot be cast to java.util.List
Caused by:
java.lang.ClassCastException: java.lang.String cannot be cast to java.util.List
at org.forgerock.openig.filter.LoggingFilter.writeHeaders(LoggingFilter.java:132)
at org.forgerock.openig.filter.LoggingFilter.logResponse(LoggingFilter.java:119)
at org.forgerock.openig.filter.LoggingFilter.filter(LoggingFilter.java:86)
at org.forgerock.openig.filter.Chain.handle(Chain.java:54)
```

This error is typically encountered when using the [AssignmentFilter](#) and setting a string value for one of the Headers. All headers are stored in Lists so the header must be addressed with a subscript. For example, if you try to set `exchange.request.headers['Location']` for a redirect in the response object, you should instead set `exchange.request.headers['Location'][0]`. A header without a subscript leads to the error above.

13.6 StaticResponseHandler results in a blank page

You must define an entity for the response. For example:

```
{
  "name": "AccessDeniedHandler",
  "type": "org.forgerock.openig.handler.StaticResponseHandler",
  "config": {
    "status": 403,
    "reason": "Forbidden",
    "entity": "<html><h2>User does not have permission</h2></html>"
  }
}
```

13.7 OpenIG is not logging users in

If you are proxying to more than one application in multiple DNS domains, you must make sure your container is enabled for domain cookies. See the chapter on [Configuring Deployment Containers](#) for details on your specific container.

Index

A

Architecture, 5

C

Configuration

- Federation, 40, 41

- HTTP & HTTPS, 56

- Login with cookie, 50

- Login with filter, 52

- Login with hidden value, 54

- Microsoft Online Outlook Web Access, 60

- Multiple applications, 57

- Proxy & capture, 47

- Simple login form, 49

Containers

- JBoss, 18

- Jetty, 17

- Supported, 15

- Tomcat, 15

Customizations

- Extension points, 63

- Filters, 63

- Handlers, 64

- Heap objects, 64

- Servlets, 65

I

Installation, 13

- Federation, 39

T

Troubleshooting, 67

Tutorials

- Basic features, 27

- Capture & relay passwords, 33

- Federation, 45

U

Use cases, 21, 22, 23, 25, 25

